

Sequential Response Surface Methodology and Central Composite Design for Multi-Objective Hyperparameter Optimization in Gradient Boosted Trees under Stochastic Nuisance Blocking

Design of Experiments & Machine Learning Research Group
Empirical Systems Engineering Laboratory

October 8, 2026

Abstract

Hyperparameter optimization (HPO) in gradient boosted decision trees is typically conducted using stochastic search methods such as uniform random sampling or Tree-structured Parzen Estimators (TPE). While effective at unconstrained objective minimization, heuristic search strategies offer limited parametric interpretability, treat stochastic nuisance variation from data partitioning as negligible noise, and lack formal statistical criteria for local curvature and model adequacy. In this investigation, we implement a Design of Experiments (DOE) framework based on the methodology of Douglas C. Montgomery. We formulate a multi-objective HPO pipeline optimizing an XGBoost regressor (100 trees) on the California Housing dataset across four factors (three continuous, one integer depth): learning rate (η), maximum tree depth, subsampling fraction, and L_2 regularization (λ). Stochastic variation is isolated via a Randomized Complete Block Design (RCBD) across five deterministic seed blocks, where the block-level nuisance component accounts for an estimated $\sigma_{\text{block}} \approx 0.0077$ RMSE in unexplained residual variance ($\text{ICC} = 0.4069, F = 20.21, p < 0.0001$). Rather than adaptive runtime stopping, the full 140-run Face-Centered Central Composite Design (FCCD) was pre-planned and executed under a single randomized execution schedule across blocks, then evaluated in sequential analytical stages: a Phase 1 screening factorial with center points ($N = 100$) revealing substantial curvature ($F = 105432.31, p < 10^{-15}$ against center pure error; $F \approx 4, 140$ against overall pure error), followed by Phase 2 axial augmentation ($N = 140$ total), decomposing the 121 df second-order residual into structural lack of fit (10 df), treatment $\times$ block interaction (96 df), and genuine center pure error (15 df). Formal lack-of-fit testing reveals significant second-order deficiency ($F = 62.21, p = 7.78e - 41$) with an RMS misfit of 0.0297. Restricting the domain ($x_1 \geq -0.5$) reduces RMS misfit to 0.0127, but leaves only two x_1 factor levels (aliasing linear and quadratic effects); significant lack of fit remains ($F = 10.11, p = 0.0001$), confirming that second-order polynomials remain local approximations across boosting loss surfaces. Residual diagnostics show consistency with normality ($W = 0.9971, p = 0.9947$) and block variance equality ($W = 0.1754, p = 0.9507$), while directional variance compression across the factor domain is accommodated via HC3 robust standard errors. Phase 3 canonical spectral analysis indicates that the unconstrained stationary point lies outside the operational domain at coded distance 17.16, with wild bootstrap analysis confirming a stationary ridge system. Constrained optimization over the operational hypercube yields a predicted optimum at depth 7 ($\hat{y} = 0.4483$). In Phase 4, Derringer-Suich multi-objective desirability optimization produces an operational compromise coordinate ($\mathbf{x}_{\text{MO}}^*$, depth 4, $\eta = 0.2324$) balancing predictive error with inference latency. In Phase 5, empirical confirmation across $m = 10$ fresh seeds confirms surrogate predictions at $\mathbf{x}_{\text{MO}}^*$ (Y_1 : 0.4853 ± 0.0115 inside Satterthwaite PI $[0.4677, 0.4932]$; Y_2 : $142.33 \pm 7.79 \mu\text{s}$ inside $[122.8, 153.4] \mu\text{s}$). At depth 7 ($\mathbf{x}_{\text{SO}}^*$), empirical validation RMSE (0.4700 ± 0.0093) exceeds the Satterthwaite prediction interval ($[0.4354, 0.4612]$) due to structural surrogate optimism ($+0.0217$), though empirical testing confirms it as a competitive configuration. Benchmark comparisons across 20 fresh seeds on a fixed external holdout test

set ($N = 4,128$) demonstrate that DOE $\mathbf{x}_{\text{MO}}^*$ delivers superior predictive accuracy at comparable latency compared to Multi-Objective TPE (0.4884 vs. 0.5032 test RMSE at $151.5 \mu\text{s}$ vs. $147.9 \mu\text{s}$, expanding point hypervolume by $+11.2\%$ from 9.88 to 10.99). Compared to single-objective TPE ($206.9 \mu\text{s}$), DOE $\mathbf{x}_{\text{MO}}^*$ reduces latency by 26.8% ($151.5 \mu\text{s}$, making TPE 36.6% slower). Empirical depth 7 ($\mathbf{x}_{\text{SO}}^*$) achieves test RMSE 0.4690, matching TPE (0.4668) with 14.5% lower latency.

Keywords: Design of Experiments (DOE), Response Surface Methodology (RSM), Central Composite Design (CCD), Gradient Boosted Trees, XGBoost, Multi-Objective Optimization, Randomized Complete Block Design (RCBD), Canonical Analysis, Model Adequacy.

1 Introduction and Theoretical Framework

Gradient Boosted Decision Trees (GBDTs), such as XGBoost [2], achieve competitive predictive accuracy on tabular data. However, tuning GBDT hyperparameters involves non-linear interactions, non-convex loss surfaces, and operational trade-offs between predictive accuracy and single-sample inference latency.

Standard practice relies on heuristic search algorithms, including uniform random search [3] and sequential Bayesian optimization methods such as Tree-structured Parzen Estimators (TPE) [4, 5] or Gaussian Process regression. In practical workflows, these black-box optimizers present three structural trade-offs relative to parametric response surface methods:

- (i) **Absence of Parametric Attribution:** Black-box optimizers treat the objective function as an opaque evaluation oracle. They cannot quantify the percentage of variance explained by main effects versus two-factor interactions, nor do they provide an analytical surrogate for sensitivity analysis.
- (ii) **Confounding of Stochastic Nuisance:** Variation arising from random data partitioning and stochastic subsampling is confounded with hyperparameter sensitivity unless structured blocking is introduced.
- (iii) **Lack of Curvature and Adequacy Tests:** Heuristic search strategies lack hypothesis tests to determine whether an operational region is planar, quadratic, or deficient in fit.

To address these limitations within an interpretable framework, this study implements a sequential Design of Experiments (DOE) pipeline based on Chapters 5, 9, 10, and 14 of Douglas C. Montgomery’s *Design and Analysis of Experiments* [1], following response surface methodology developed by Box and Wilson [7] and extended by Myers et al. [8]. By combining randomized complete block screening, curvature testing, Central Composite Design (CCD) augmentation, spectral canonical analysis, and Derringer-Suich desirability functions [6], we evaluate the extent to which classical response surface methodology can optimize GBDTs.

All metrics, sums of squares, test statistics, and diagnostics reported in this paper are derived from model fits on the California Housing benchmark ($N = 20,640$). To prevent test data leakage, an external holdout test set of 4,128 observations (20%) was set aside prior to experimental design and model training. The remaining 16,512 observations formed the development pool, which each seed block partitioned into 12,384 training (75%) and 4,128 validation (25%) samples. While holdout test metrics were concurrently recorded in the execution log for subsequent verification, the holdout test data remained strictly sequestered from all surrogate fitting, curvature testing, lack-of-fit analysis, and desirability optimization.

2 Experimental Architecture and Mathematical Formulation

2.1 Operational Factors and Standard Coding

We define four hyperparameter factors (three continuous, one integer depth) controlling structural complexity and regularization for an XGBoost regressor with $n_{\text{estimators}} = 100$:

- **Factor A** (x_1): Learning rate ($\eta \in [0.01, 0.30]$), coded on a logarithmic scale: $\xi_1 = \ln(\eta) \in [\ln(0.01), \ln(0.30)] = [-4.6052, -1.2040]$.
- **Factor B** (x_2): Maximum tree depth (depth $\in [3, 9]$), controlling structural model complexity.
- **Factor C** (x_3): Subsample fraction (subsample $\in [0.50, 1.00]$), introducing stochastic row subsampling.
- **Factor D** (x_4): L_2 leaf regularization ($\lambda \in [0.10, 10.00]$), coded on a logarithmic scale: $\xi_4 = \ln(\lambda) \in [\ln(0.10), \ln(10.00)] = [-2.3026, 2.3026]$.

Each natural variable ξ_i is mapped to dimensionless coded space $x_i \in [-1, +1]$ via:

$$x_i = \frac{\xi_i - \frac{\xi_{i,\max} + \xi_{i,\min}}{2}}{\frac{\xi_{i,\max} - \xi_{i,\min}}{2}}, \quad i \in \{1, 2, 3, 4\} \quad (1)$$

The operational domain forms a 4-dimensional hypercube $\mathcal{D} = [-1, +1]^4$. Table 1 summarizes factor coding.

Factor	Natural Variable	Transform	Low −1	Center 0	High +1
Factor A (x_1)	Learning Rate (η)	$\ln(\eta)$	0.0100	0.0548	0.3000
Factor B (x_2)	Max Tree Depth	Linear	3	6	9
Factor C (x_3)	Subsample Fraction	Linear	0.5000	0.7500	1.0000
Factor D (x_4)	L_2 Reg (λ)	$\ln(\lambda)$	0.1000	1.0000	10.0000

Table 1: Hyperparameter Factors, Operational Domains, and Design Levels

2.2 Randomized Complete Block Design (RCBD) for Stochastic Nuisance

Stochastic variation in data partitioning and subsampling introduces non-trivial variance across evaluation trials. We structure the experimental protocol as a Randomized Complete Block Design (RCBD) across $b = 5$ distinct seed blocks: $\mathcal{S} = \{42, 101, 202, 303, 404\}$. Table 2 summarizes the assignment of split seeds and model seeds across phases.

Phase / Design Point	Point IDs	Block	Split Seed	Model Seed	Reps	Isolated Variance
Phase 1: Corner Runs	1...16	$b \in \{1..5\}$	S_b	S_b	1	Factorial + Block
Phase 1: Center Points	17	$b \in \{1..5\}$	S_b	$S_b + 1000r$	4	Subsampling Error (15 df)
Phase 2: Axial Runs	18...25	$b \in \{1..5\}$	S_b	S_b	1	Quadratic + Block

Table 2: Structure of Blocks, Split Seeds, and Model Seeds Across Experimental Phases

Specifically, each block b defines a distinct 75/25 train/validation partition of the 16,512 development observations via split seed S_b . For corner runs (1...16) and axial runs (18...25), the internal XGBoost subsampling seed is also set to S_b , meaning that the block factor captures the joint nuisance impact of data-split selection and stochastic subsampling realization. For the four center point replicates within block b , the train/validation split is held fixed at S_b while the model subsampling seed is varied via $S_b + r \times 1000$ ($r \in \{0, 1, 2, 3\}$), isolating within-block stochastic subsampling pure error (15 degrees of freedom) from block-to-block variation (4 degrees of freedom).

The statistical model for response y_{uj} at design point u in block j is:

$$y_{uj} = \mu + \tau_u + \beta_j + \epsilon_{uj}, \quad j \in \{1, \dots, 5\} \quad (2)$$

where τ_u represents the treatment effect of design coordinate $\mathbf{x}_u$, β_j represents the block effect, and ϵ_{uj} is experimental error.

To quantify the magnitude of variance absorbed by blocking, we compute the Intraclass Correlation Coefficient (ICC) under both ANOVA and REML formulations:

$$\text{ICC} = \frac{\sigma_{\text{block}}^2}{\sigma_{\text{block}}^2 + \sigma_{\epsilon}^2} \quad (3)$$

For validation RMSE, the ANOVA estimator yields $\text{ICC} = 0.4069$ ($F = 20.21, p < 0.0001$), which matches the restricted maximum likelihood estimate ($\text{ICC}_{\text{REML}} = 0.4069$). Importantly, this 40.7% represents a share of *unexplained residual variance* rather than total experimental variation (where primary factors account for $> 90\%$ of total sum of squares). In absolute terms, the estimated block standard deviation is $\sigma_{\text{block}} = \sqrt{\sigma_{\text{block}}^2} \approx 0.0077$ RMSE. This nuisance variation is comparable in magnitude to the performance differences separating competitive hyperparameter configurations (often 0.005–0.015 RMSE), which makes blocking necessary when comparing close configurations.

While the block effect is non-significant in Phase 1 ($F = 0.33, p = 0.8541$), it is significant in the Phase 2 second-order model ($F = 20.21, p < 10^{-11}$). This difference reflects variance shielding: in Phase 1, unmodeled quadratic curvature pooled into the residual denominator ($\text{MS}_E = 0.003485$), inflating error variance and attenuating the block F -ratio. Once second-order terms accounted for curvature in Phase 2 (MS_E reduced to 0.000086), seed partition variance became detectable. For inference latency, the block effect is negligible ($\text{ICC} = 0.0000, F = 0.75, p = 0.5591$), as execution speed is governed by hardware architecture and tree topology rather than data splits.

2.3 Dual Responses and Latency Measurement Protocol

For every trial, two responses are measured:

1. **Response Y_1 (Validation RMSE):** Root mean squared error evaluated on the 25% validation partition ($N_{\text{val}} = 4,128$) of the development pool. The holdout 20% test partition ($N_{\text{test}} = 4,128$) remains strictly untouched until final generalization testing.
2. **Response Y_2 (Single-Sample Inference Latency):** Execution time in microseconds per single sample ($\mu\text{s}/\text{sample}$) on CPU. We emphasize that this represents single-sample point evaluation rather than batched throughput, and the quadratic depth term is an empirical response-surface fit across sampled depths $d \in \{3, 6, 9\}$.

To eliminate session drift, CPU frequency fluctuations, and OS scheduling noise, all latency measurements follow a unified measurement protocol:

- (i) **Pinned CPU Affinity:** The execution process is pinned to CPU Core 0 via Windows API `SetProcessAffinityMask(GetCurrentProcess(), 1)`.
- (ii) **Warmup and Single-Threaded Execution:** 100 untimed warmup iterations (50 per prediction interface) precede timed execution, with XGBoost configured for single-threaded inference (`n_jobs=1`).
- (iii) **Repeated Timed Evaluations:** Inference latency is timed across 1,000 single-sample evaluations using high-resolution monotonic clocks (`time.perf_counter_ns`), recording both the scikit-learn `predict` method and native `inplace_predict`. In XGBoost 3.2.0, `predict` on 2D arrays already executes in-place internally, so the measured 39–47 μs divergence reflects Python and scikit-learn wrapper overhead and array validation rather than memory copying.
- (iv) **Interleaved Benchmark Runs:** During benchmarking, evaluation runs across competing algorithms are interleaved within the same pinned process session to maintain consistent CPU thermal and operating states.

3 Phase 1: Screening and Curvature Assessment

Phase 1 constructs a 2^4 full factorial design consisting of $n_F = 16$ corner runs, augmented with $n_C = 4$ center points per block ($x_1 = x_2 = x_3 = x_4 = 0$). Evaluated across all 5 seed blocks, Phase 1 comprises $N_1 = 100$ genuine training runs.

We fit a first-order response surface model with all two-factor interaction terms:

$$Y = \beta_0 + \sum_{i=1}^4 \beta_i x_i + \sum_{i < j} \beta_{ij} x_i x_j + \sum_{b=2}^5 \gamma_b Z_b + \epsilon \quad (4)$$

where Z_b are dummy indicators for blocks 2 through 5. Table 3 details the Type III Sum of Squares ANOVA for Response Y_1 .

Source of Variation	SS	DF	MS	F-Stat	p value	Partial η^2
Intercept	7.309127	1	7.309127	2097.15	$< 10^{-15}$	0.9610
Block Effect $C(\text{block})$	0.004662	4	0.001166	0.33	0.8541	0.0155
Factor A: x_1 ($\ln(\eta)$)	1.245625	1	1.245625	357.40	$< 10^{-15}$	0.8079
Factor B: x_2 (Depth)	0.082155	1	0.082155	23.57	0.0000	0.2171
Factor C: x_3 (Subsample)	0.002406	1	0.002406	0.69	0.4084	0.0081
Factor D: x_4 ($\ln(\lambda)$)	0.000089	1	0.000089	0.03	0.8731	0.0003
Interaction $x_1 \cdot x_2$	0.091725	1	0.091725	26.32	0.0000	0.2364
Interaction $x_1 \cdot x_3$	0.001878	1	0.001878	0.54	0.4649	0.0063
Interaction $x_1 \cdot x_4$	0.007452	1	0.007452	2.14	0.1474	0.0245
Interaction $x_2 \cdot x_3$	0.001106	1	0.001106	0.32	0.5747	0.0037
Interaction $x_2 \cdot x_4$	0.000092	1	0.000092	0.03	0.8714	0.0003
Interaction $x_3 \cdot x_4$	0.000081	1	0.000081	0.02	0.8790	0.0003
Residual Error	0.296247	85	0.003485	—	—	—

Table 3: Phase 1 ANOVA Table (Type III Sum of Squares) for Validation RMSE (Y_1)

Factor A ($\ln(\eta)$) dominates predictive performance ($F = 357.40$, partial $\eta^2 = 0.8079$), followed by Factor B (tree depth, $F = 23.57$) and the interaction $x_1 \cdot x_2$ ($F = 26.32$, $p < 10^{-6}$).

3.1 Formal Curvature Hypothesis Test

Following Montgomery (Chapter 10), curvature is tested by comparing the average response of factorial points ($\bar{y}_F$) against center points ($\bar{y}_C$):

$$SS_{\text{Curvature}} = \frac{n_F n_C (\bar{y}_F - \bar{y}_C)^2}{n_F + n_C} \quad (5)$$

Across the 100 runs evaluated in Phase 1:

$$n_F = 80, \quad \bar{y}_F = 0.6333, \quad n_C = 20, \quad \bar{y}_C = 0.4989, \quad \bar{y}_F - \bar{y}_C = 0.1343 \quad (6)$$

yielding $SS_{\text{Curvature}} = 0.288786$. Estimating pure error variance from center point replicates within blocks (subsampling pure error, $df = 15$) yields $MS_{\text{PE, Center}} = 2.74e - 06$. The resulting test statistic is:

$$F_{\text{Curvature}} = \frac{MS_{\text{Curvature}}}{MS_{\text{PE, Center}}} = 105432.31 \quad (p < 10^{-15}) \quad (7)$$

If evaluated against overall pure error pooled across all replicated design coordinates ($df = 83$, $MS_{\text{PE, All}} = 6.98 \times 10^{-5}$), the F -statistic is also significant ($F \approx 4,140$, $p < 10^{-15}$). Under either reference, the null hypothesis of planarity is rejected, confirming curvature that requires second-order augmentation.

4 Phase 2: Central Composite Design Augmentation

4.1 Face-Centered Central Composite Design (FCCD)

To fit a second-order polynomial without discarding Phase 1 runs, the design is augmented into a Face-Centered Central Composite Design (FCCD, $\alpha = 1.0$), adding $2k = 8$ axial runs evaluated across the 5 seed blocks (40 additional runs, cumulative $N = 140$).

The second-order response surface model is:

$$Y = \beta_0 + \sum_{i=1}^4 \beta_i x_i + \sum_{i=1}^4 \beta_{ii} x_i^2 + \sum_{i < j} \beta_{ij} x_i x_j + \sum_{b=2}^5 \gamma_b Z_b + \epsilon \quad (8)$$

Table 4 reports the second-order ANOVA for Validation RMSE ($R^2 = 0.9951$, Adj $R^2 = 0.9944$). Because residual variance varies across the factor domain, both OLS and HC3 heteroscedasticity-consistent standard errors and p-values are reported.

Source of Variation	SS	DF	MS	F-Stat	OLS p	HC3 SE	HC3 p
Intercept	4.718852	1	4.718852	54749.42	$< 10^{-15}$	0.0021	$< 10^{-15}$
Block Effect $C(\text{block})$	0.006966	4	0.001742	20.21	< 0.0001	—	—
Factor A: x_1 ($\ln(\eta)$)	1.388673	1	1.388673	16111.76	$< 10^{-15}$	0.0012	$< 10^{-15}$
Factor B: x_2 (Depth)	0.101706	1	0.101706	1180.02	$< 10^{-15}$	0.0013	$< 10^{-15}$
Factor C: x_3 (Subsample)	0.002013	1	0.002013	23.36	< 0.0001	0.0012	0.0001
Factor D: x_4 ($\ln(\lambda)$)	0.000044	1	0.000044	0.52	0.4743	0.0012	0.5483
Quadratic x_1^2	0.146265	1	0.146265	1697.01	$< 10^{-15}$	0.0023	$< 10^{-15}$
Quadratic x_2^2	0.008779	1	0.008779	101.85	$< 10^{-15}$	0.0035	< 0.0001
Quadratic x_3^2	0.000005	1	0.000005	0.05	0.8188	0.0023	0.8005
Quadratic x_4^2	0.000006	1	0.000006	0.08	0.7844	0.0020	0.7197
Interaction $x_1 \cdot x_2$	0.091725	1	0.091725	1064.22	$< 10^{-15}$	0.0013	$< 10^{-15}$
Interaction $x_1 \cdot x_3$	0.001878	1	0.001878	21.79	< 0.0001	0.0013	0.0003
Interaction $x_1 \cdot x_4$	0.007452	1	0.007452	86.46	$< 10^{-15}$	0.0013	< 0.0001
Interaction $x_2 \cdot x_3$	0.001106	1	0.001106	12.83	0.0005	0.0013	0.0050
Interaction $x_2 \cdot x_4$	0.000092	1	0.000092	1.07	0.3040	0.0013	0.4118
Interaction $x_3 \cdot x_4$	0.000081	1	0.000081	0.94	0.3334	0.0013	0.4400
Residual Error	0.010429	121	0.000086	—	—	—	—

Table 4: Full Second-Order ANOVA Table for Validation RMSE (Y_1) with HC3 Robust Standard Errors

The ANOVA confirms substantial quadratic curvature along x_1^2 ($F = 1,697.01, p < 10^{-15}$; HC3 $p < 10^{-15}$) and x_2^2 ($F = 101.85, p < 10^{-15}$; HC3 $p < 0.0001$). Under HC3 covariance estimation, structural effects $x_1, x_2, x_1^2, x_2^2, x_1 x_2$, and $x_1 x_4$ remain significant ($p < 0.0001$). Table 5 provides the corresponding ANOVA for Inference Latency (Y_2).

Source of Variation	SS	DF	MS	F-Stat	p value
Intercept	555450.04	1	555450.04	1903.08	$< 10^{-15}$
Block Effect $C(\text{block})$	876.94	4	219.24	0.75	0.5591
Factor A: x_1 ($\ln(\eta)$)	1265.46	1	1265.46	4.34	0.0394
Factor B: x_2 (Depth)	236841.33	1	236841.33	811.46	$< 10^{-15}$
Factor C: x_3 (Subsample)	179.04	1	179.04	0.61	0.4350
Factor D: x_4 ($\ln(\lambda)$)	483.26	1	483.26	1.66	0.2006
Quadratic x_1^2	1576.98	1	1576.98	5.40	0.0218
Quadratic x_2^2	11634.52	1	11634.52	39.86	< 0.0001
Quadratic x_3^2	8.35	1	8.35	0.03	0.8660
Quadratic x_4^2	116.71	1	116.71	0.40	0.5284
Interaction $x_1 \cdot x_2$	662.90	1	662.90	2.27	0.1344
Interaction $x_1 \cdot x_3$	14.00	1	14.00	0.05	0.8270
Interaction $x_1 \cdot x_4$	829.24	1	829.24	2.84	0.0945
Interaction $x_2 \cdot x_3$	813.68	1	813.68	2.79	0.0976
Interaction $x_2 \cdot x_4$	47.57	1	47.57	0.16	0.6871
Interaction $x_3 \cdot x_4$	52.53	1	52.53	0.18	0.6721
Residual Error	35316.19	121	291.87	—	—

Table 5: Full Second-Order ANOVA Table for Inference Latency (Y_2 , $\mu\text{s}/\text{sample}$)

4.2 Lack-of-Fit Decomposition and Structural Model Adequacy

In response surface methodology [1], evaluating model adequacy requires decomposing the model residual sum of squares into pure experimental error and structural lack of fit. In our blocked central composite design ($N = 140$ runs across 5 blocks), non-center design coordinates are observed once per block. Genuine within-coordinate and within-block replication exists exclusively at the center points ($x_1 = x_2 = x_3 = x_4 = 0$), where 4 model-seed replicates within each of the 5 blocks provide exactly $\text{df}_{\text{PE, center}} = 5 \times (4 - 1) = 15$ degrees of freedom of genuine within-cell pure replicate error ($\text{SS}_{\text{PE, center}} = 0.000041$, $\text{MS}_{\text{PE, center}} = 2.74e - 06$).

Fitting a saturated additive cell model containing all $m = 25$ unique design coordinates and 4 block parameters (29 parameters) leaves $140 - 29 = 111$ degrees of freedom with $\text{SS} = 0.001579$ ($\text{MS} = 1.42e - 05$). Rather than representing uniform replicate noise, these 111 df decompose into:

- (i) Genuine center-point pure replicate error (15 df, $\text{SS} = 0.000041$);
- (ii) Treatment $\times$ Block interaction / split-specific treatment response ($111 - 15 = 96$ df, $\text{SS} = 0.001538$, $\text{MS} = 1.60e - 05$).

Evaluating the interaction against center pure error yields $F = 5.85$ ($p = 0.0002$), establishing statistically significant split-by-hyperparameter interaction: different train/validation data splits exhibit subtle variations in optimal response curvature.

The remaining $121 - 111 = 10$ df represent structural lack of fit ($\text{SS}_{\text{LoF}} = 0.008850$, $\text{MS}_{\text{LoF}} = 0.000885$). Table 6 details this complete decomposition for both responses.

Source of Variation / Model	SS	DF	MS	F	p value	Reference / RMS Misfit
Decomposition of Second-Order Model Residual (Y_1: Validation RMSE, Full CCD)						
Structural Lack of Fit	0.008850	10	0.000885	323.10	$< 10^{-15}$	vs. Center PE
Treatment $\times$ Block Interaction	0.001538	96	1.60×10^{-5}	5.85	0.0002	vs. Center PE
Genuine Center Pure Error	0.000041	15	2.74×10^{-6}	—	—	Base Replicate
Total Model Residual	0.010429	121	8.62×10^{-5}	—	—	RMS = 0.0093
Saturated Additive Baseline	0.001579	111	1.42×10^{-5}	62.21*	$< 10^{-15}$	(*LoF vs. Additive)
Restricted Domain Sensitivity (Y_1: $x_1 \geq -0.5$, $\eta \geq 0.033$, $N = 95$)						
Structural Lack of Fit	0.000324	2	0.000162	10.11	0.0001	RMS = 0.0127
Additive Baseline Residual	0.001203	75	1.60×10^{-5}	—	—	Saturated Base
Total Restricted Residual	0.001528	77	1.98×10^{-5}	—	—	RMS = 0.0045
Latency Residual Decomposition and Adequacy (Y_2: Single-Sample Latency, μs)						
Structural Lack of Fit	3322.37	10	332.24	1.20	0.3625	vs. Center PE (RMS = 18.23)
Treatment $\times$ Block Interaction	27844.93	96	290.05	1.05	0.4906	vs. Center PE
Genuine Center Pure Error	4148.89	15	276.59	—	—	Base Replicate
Total Model Residual	35316.19	121	291.87	1.15*	0.3305	(*LoF vs. Additive)

Table 6: Residual Variance Decomposition (121 df), Lack-of-Fit Tests, and Domain Restriction Sensitivity

For Validation RMSE (Y_1) on the full operational domain:

- Tested against center pure error (15 df), lack of fit is statistically significant ($F = 323.10, p < 10^{-15}$).
- Tested against the saturated additive baseline (111 df, pooling interaction and replicate noise), lack of fit is also significant ($F = 62.21, p = 7.78e - 41$).

Under either reference baseline, the root-mean-square misfit per degree of freedom is $\sqrt{MS_{LoF}} = 0.0297$, which is approximately $7.8\times$ the additive residual standard deviation ($\sqrt{MS_{PE}} = 0.0038$) and $17.5\times$ the center pure error standard deviation ($\sqrt{MS_{PE, center}} = 0.0017$). By contrast, the unweighted root-mean-square cell prediction error across the 25 unique design points is $\sqrt{\frac{1}{25} \sum_{i=1}^{25} (\bar{y}_i - \hat{y}_i)^2} \approx 0.0084$, while the observation-weighted root-mean-square misfit across all 140 design evaluations is $\sqrt{SS_{LoF}/140} \approx 0.0080$ (down-weighting non-center cells relative to the replicated center).

To inspect where this lack of fit concentrates across the hypercube, per-design-point residual contributions $n_i(\bar{y}_i - \hat{y}_i)^2$ reveal that the single largest lack-of-fit discrepancy occurs at corner $(1, 1, -1, -1)$ ($x_1 = 1, x_2 = 1, x_3 = -1, x_4 = -1$, contribution ≈ 0.00153), where high learning rate and depth combined with low subsampling and regularization challenge second-order surface adequacy. Substantial discrepancies also concentrate at low learning rates ($x_1 = -1, \eta = 0.01$), notably at $(-1, -1, 1, 1)$ (0.00097) and $(-1, 1, -1, 1)$ (0.00071), reflecting underfitting when 100 boosting iterations are paired with low step sizes. When refitting on the restricted sub-domain $x_1 \geq -0.5$ ($\eta \geq 0.033$), the sum of squares of lack of fit drops by 96.3% (from 0.008850 to 0.000324), and the RMS misfit declines by 57% to 0.0127. However, we emphasize that dropping $x_1 = -1$ leaves only two factor levels ($x_1 \in \{0, 1\}$). Because two levels induce the identity $x_1^2 \equiv x_1$, separate linear and quadratic coefficients along Factor A are unidentifiable from those levels alone, reducing df_{LoF} from 10 to 2. Consequently, this restricted test cannot assess quadratic curvature across $x_1 \in [0, 1]$ where all candidate optima reside. Even within this restricted region, lack of fit remains statistically significant ($F = 10.11, p = 0.0001$), with RMS

misfit (0.0127) remaining roughly $3\times$ the additive residual standard deviation (0.0040). This confirms that while low- η underfitting drives the largest discrepancies, a second-order polynomial remains an approximation with non-negligible structural bias across gradient boosting loss surfaces.

For Latency (Y_2), structural lack of fit is non-significant against both reference baselines ($F = 1.20, p = 0.3625$ against center pure error; $F = 1.15, p = 0.3305$ against the saturated additive baseline). The treatment $\times$ block interaction is also non-significant ($F = 1.05, p = 0.4906$), confirming that latency is invariant across seed blocks and data splits. To characterize the functional form of latency growth, five candidate models were compared against empirical measurements, as reported in Table 7. For the log-linear specification, AIC and BIC values are adjusted via the Jacobian of transformation ($+2 \sum \ln(y_i) = 1,450.8$) to allow direct likelihood comparison on the natural microsecond scale.

Candidate Latency Model	R^2	Adj R^2	AIC (Adjusted)	BIC (Adjusted)	RMSE (μs)
Linear (depth)	0.792	0.790	1255.2	1261.1	21.26
Quadratic (depth + depth ²)	0.859	0.857	1202.6	1211.4	17.56
Log-Linear ($\ln(\text{latency})$ depth)	0.817	0.815	1218.1	1224.0	19.57
Multi-Factor (depth, depth ² , eta, sub, lam)	0.868	0.863	1199.5	1217.1	17.18
Full CCD Second-Order (All factors + Block)	0.882	0.864	1209.6	1265.5	17.08

Table 7: Comparison of Candidate Regression Models for Inference Latency (Y_2)

The quadratic latency model ($R^2 = 0.859$, adjusted AIC = 1202.6) yields lower AIC than the linear model ($R^2 = 0.792$, AIC = 1255.2) and the log-linear model (adjusted AIC = 1218.1), consistent with the significant x_2^2 term in ANOVA ($F = 39.86, p < 0.0001$). Latency exhibits super-linear growth across the empirical depth range ($d \in \{3, 6, 9\}$), with the quadratic fit capturing accelerating tree traversal costs while avoiding the severe over-prediction of exponential formulations. Comparing standard prediction against `inplace_predict` shows an execution gap of 39–47 μs , indicating that Python and scikit-learn wrapper overhead adds fixed baseline latency while base traversal scales with tree depth.

5 Statistical Residual Diagnostics

To evaluate regression assumptions, diagnostic checks were conducted on externally studentized residuals:

- (a) **Normality:** Evaluated via Shapiro-Wilk on studentized residuals: $W = 0.9971, p = 0.9947$. The test fails to reject the null hypothesis of normality at $\alpha = 0.05$.
- (b) **Variance Homogeneity across Blocks:** Assessed via Levene’s test across the 5 seed blocks: $W = 0.1754, p = 0.9507$ (fails to reject equal block variance). Brown-Forsythe testing across the 25 design points yields $W = 0.6872, p = 0.8552$, failing to reject uniform replicate dispersion across design coordinates.
- (c) **Reconciliation of Breusch-Pagan and Brown-Forsythe:** The Breusch-Pagan test detects continuous heteroscedasticity across the four-factor hypercube (LM = 70.10, $p < 0.0001$). This does not contradict the Brown-Forsythe block homogeneity: while variance across seed blocks is uniform, residual variance compresses directionally toward regions

of higher η and deeper trees where the boosting loss reaches an asymptote. This heteroscedasticity motivates using HC3 robust standard errors for coefficient tests in Table 4; note, however, that HC3 standard errors do not correct for non-constant variance when constructing ordinary prediction intervals.

- (d) **Independence and Run Order:** Durbin-Watson statistic evaluated in actual randomized execution order yields $DW = 1.9918$. The Ljung-Box test ($Q = 4.01, p = 0.5486$) and Wald-Wolfowitz Runs test ($p = 0.8576$) fail to detect serial correlation or thermal execution drift.
- (e) **Influence and Leverage:** Maximum Cook's distance is $D_{\max} = 0.097$, well below the commonly cited heuristic reference cutoff of 1.0.

Figure 1 displays the 4-in-1 residual diagnostic panel.




Figure 1: Residual Diagnostics Panel for Second-Order Validation RMSE Model: (a) Studentized residuals vs. fitted values; (b) Normal Q-Q probability plot ($p = 0.9947$); (c) Residuals vs. randomized run execution order ($DW = 1.9918$, Ljung-Box $p = 0.5486$, Runs test $p = 0.8576$); (d) Cook's distance influence plot ($D_{\max} = 0.097$, reference guideline 1.0).

6 Phase 3: Canonical and Ridge Analysis

6.1 Matrix Representation and Stationary Point

The fitted second-order model is expressed in matrix notation:

$$\hat{y} = b_0 + \mathbf{x}^T \mathbf{b} + \mathbf{x}^T \mathbf{B} \mathbf{x} \quad (9)$$

where $b_0 = 0.4994$ (averaging block effects equally), $\mathbf{b} = [-0.1242, -0.0336, -0.0047, -0.0007]^T$, and:

$$\mathbf{B} = \begin{bmatrix} 0.1065 & 0.0169 & -0.0024 & -0.0048 \\ 0.0169 & 0.0261 & -0.0019 & -0.0005 \\ -0.0024 & -0.0019 & 0.0006 & 0.0005 \\ -0.0048 & -0.0005 & 0.0005 & 0.0007 \end{bmatrix} \quad (10)$$

The unconstrained stationary point $\mathbf{x}_0$ satisfies $\mathbf{b} + 2\mathbf{B}\mathbf{x}_0 = \mathbf{0}$:

$$\mathbf{x}_0 = -\frac{1}{2}\mathbf{B}^{-1}\mathbf{b} = [0.4171, 1.3524, 15.6944, -6.8009]^T \quad (11)$$

The stationary coordinate along Factor C ($x_3 = 15.6944$) lies far outside the operational cube $[-1, +1]^4$. Its Euclidean distance from the coded origin is $\|\mathbf{x}_0\| = 17.16$, while its Euclidean distance from the boundary of the feasible hypercube is 15.80 coded units. The unconstrained stationary point has a predicted response of $\hat{y}_0 = 0.4161$, an extrapolation artifact indicating that unconstrained quadratic optimization yields an infeasible point.

6.2 Spectral Decomposition and Ridge Analysis

Eigenvalues of $\mathbf{B}$ are $\lambda_1 = 0.000057, \lambda_2 = 0.000873, \lambda_3 = 0.022759, \lambda_4 = 0.110196$ ($\sum \lambda_i = \text{trace}(\mathbf{B}) = 0.133886$). While point estimates are positive, Rademacher wild bootstrap (2,000 replications) reveals that the 95% bootstrap interval for λ_1 spans $[-0.0039, 0.0022]$, with 68.2% of bootstrap samples yielding $\min(\lambda) \leq 0$. The canonical eigenvectors (columns of the orthogonal decomposition matrix $\mathbf{M}$) show that the smallest eigenvalue λ_1 aligns along $\mathbf{v}_1 = [0.023, -0.053, -0.724, 0.687]^T$, representing a contrast between Factor C (subsampling, -0.724) and Factor D (L_2 regularization, $+0.687$), while λ_2 corresponds to $\mathbf{v}_2 = [-0.043, -0.037, -0.687, -0.725]^T$, representing a joint direction along subsampling and regularization. By contrast, the dominant quadratic curvature axes align along Factor B (depth, $\mathbf{v}_3 = [-0.198, 0.978, -0.060, 0.018]^T$, $\lambda_3 = 0.022759$) and Factor A (learning rate, $\mathbf{v}_4 = [-0.979, -0.198, 0.025, 0.044]^T$, $\lambda_4 = 0.110196$). The system is classified as a **stationary or rising ridge**, with curvature along x_3 and x_4 indistinguishable from zero.

Constrained optimization over the hypercube $[-1, 1]^4$ with integer depth $d \in \{3, \dots, 9\}$ yields the results in Table 8.

Depth	Coded x_1	Coded x_2	Coded x_3	Coded x_4	Optimal η	Predicted RMSE	SE ($\hat{y}$)
3	0.8103	-1.0000	1.00	1.00	0.2173	0.4909	0.0032
4	0.7573	-0.6667	1.00	1.00	0.1985	0.4724	0.0030
5	0.7043	-0.3333	1.00	1.00	0.1814	0.4592	0.0032
6	0.6513	+0.0000	1.00	1.00	0.1658	0.4511	0.0033
7	0.5983	+0.3333	1.00	1.00	0.1515	0.4483	0.0032
8	0.5453	+0.6667	1.00	1.00	0.1384	0.4506	0.0031
9	0.4923	+1.0000	1.00	1.00	0.1265	0.4582	0.0034

Table 8: Best Predicted Validation RMSE at Each Discrete Tree Depth

The constrained single-objective optimum within the hypercube occurs at depth 7 ($\mathbf{x}_{\text{SO}}^*$: $x_1 = 0.5983, x_2 = 0.3333, x_3 = 1.0000, x_4 = 1.0000$) with predicted RMSE 0.4483. Figures 2 and 3 display the response surfaces.




Figure 2: Response Surface Topography for Validation RMSE (Y_1): (a) 2D contour slice across Factor A ($\ln(\eta)$) and Factor B (depth), passing through the constrained optimum ($x_3 = 1, x_4 = 1$). Arrow indicates trajectory toward unconstrained stationary point outside domain; (b) 3D wireframe projection.




Figure 3: Inference Latency Modelling: (a) Empirical mean latency with 95% CI error bars vs. linear ($R^2 = 0.792$) and quadratic ($R^2 = 0.859$) fits. Note the linear base intercept ($\approx 80.8 \mu s$) arising from Python execution overhead; (b) Distribution boxplots across discrete depths.

7 Phase 4: Multi-Objective Desirability Optimization

Minimizing predictive error requires deeper trees, which increases inference latency. To balance predictive error against inference latency, we use Derringer-Suich desirability functions [6]:

$$d_r(Y_r) = \begin{cases} 1.0, & Y_r < L_r \\ \left(\frac{U_r - Y_r}{U_r - L_r} \right)^{s_r}, & L_r \leq Y_r \leq U_r \\ 0.0, & Y_r > U_r \end{cases} \quad (12)$$

The composite desirability D is formulated with explicit importance weights (w_1, w_2):

$$D = (d_1(Y_1)^{w_1} \cdot d_2(Y_2)^{w_2})^{\frac{1}{w_1 + w_2}} \quad (13)$$

In the baseline scenario, equal weights $w_1 = 1, w_2 = 1$ are used with engineering limits $L_1 = 0.450, U_1 = 0.700$ for RMSE and $L_2 = 100.0, U_2 = 180.0 \mu s$ for latency ($s_1 = s_2 = 1.0$).

To select a deployable operating point, desirability D was evaluated across a high-resolution factor grid ($21 \times 7 \times 2 \times 2 = 588$ design coordinates) with discrete depths $d \in \{3, \dots, 9\}$. While unconstrained maximization pushes toward the upper boundary $\eta \rightarrow 0.30$ ($x_1 = 1.0$, reaching $D = 0.6952$) due to the 100-tree ceiling, operating directly on boundary edges risks boundary extrapolation. Holding depth $d = 4$, subsample 1.0, and $\lambda = 0.830$, composite desirability varies smoothly across the interior: $D = 0.6662$ at $\eta = 0.20$, $D = 0.6782$ at $\eta = 0.2324$, $D = 0.6833$ at $\eta = 0.25$, and $D = 0.6943$ at $\eta = 0.30$. We chose the interior coordinate at $x_1 = 0.8500$ ($\eta = 0.2324$) as the selected operational multi-objective compromise configuration $\mathbf{x}_{\text{MO}}^*$:

$$\begin{aligned} \mathbf{x}_{\text{MO}}^* &= [0.8500, -0.6667, 1.0000, -0.0812]^T \\ \implies (\eta = 0.2324, \text{depth} = 4, \text{subsample} = 1.0000, \lambda = 0.8295) \end{aligned} \quad (14)$$

achieving $D = 0.6782$ ($d_1 = 0.8783, d_2 = 0.5236$). Predicted metrics at $\mathbf{x}_{\text{MO}}^*$ are $\hat{Y}_1 = 0.4804$ and $\hat{Y}_2 = 138.11 \mu s$.

Table 9 reports sensitivity across alternative specification limits and weights. Across all scenarios, the optimal depth clusters between 3 and 5, indicating that shallow-to-moderate tree depths provide the best balance between accuracy and latency.

Scenario	RMSE Bounds [L_1, U_1]	Latency Bounds [L_2, U_2]	Weights (w_1, w_2)	Opt Depth	Opt η	Best D
Standard Config	[0.45, 0.70]	[100, 180]	(1, 1)	4	0.3000	0.6952
Strict Latency	[0.45, 0.70]	[90, 140]	(1, 2)	3	0.3000	0.2206
Strict Accuracy	[0.44, 0.60]	[100, 200]	(2, 1)	5	0.2135	0.7293
Narrow Range	[0.46, 0.65]	[110, 160]	(1, 1)	4	0.3000	0.6597

Table 9: Desirability Optimization Sensitivity Analysis Across Bounds and Weights

Figure 4 illustrates the empirical Pareto frontier and the locations of $\mathbf{x}_{\text{MO}}^*$ and $\mathbf{x}_{\text{SO}}^*$.




Figure 4: Multi-Objective Pareto Trade-off between Validation RMSE and Inference Latency. Scatter points depict individual run replicates across the 5 seed blocks, colored by tree depth using `cividis`. The empirical Pareto frontier is traced in red connecting non-dominated individual design run realizations across all blocks (constructed directly from individual evaluations in `plots.py`), with selected DOE multi-objective compromise $\mathbf{x}_{\text{MO}}^*$ (depth 4), DOE single-objective candidate $\mathbf{x}_{\text{SO}}^*$ (depth 7), confirmation means ($m = 10$), and benchmark incumbents highlighted.

8 Phase 5: Empirical Confirmation and Benchmark Comparison

8.1 Confirmation Experiments at Multi-Objective and Single-Objective Optima

To test surrogate fidelity, $m = 10$ confirmation trials were executed at $\mathbf{x}_{\text{MO}}^*$ across fresh seeds disjoint from the design ($\mathcal{S}_{\text{conf}} = \{505, \dots, 1414\}$). Because fresh seeds introduce block variance (σ_{block}^2 estimated with only 4 df) alongside model parameter covariance (MS_E with 121 df), treating block variance uncertainty as negligible would understate interval width. The total theoretical variance for the confirmation mean $\bar{y}_m$ at coordinate $\mathbf{x}_0$ with leverage $h_0 = \mathbf{x}_0^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{x}_0$ is:

$$V_{\text{conf}} = \sigma_E^2 \left(\frac{1}{m} + h_0 \right) + \sigma_{\text{block}}^2 \left(\frac{1}{m} + \frac{1}{5} \right) \quad (15)$$

Under the randomized block design with $a = 28$ runs per block, the expected mean squares satisfy $\mathbb{E}[\text{MS}_E] = \sigma_E^2$ and $\mathbb{E}[\text{MS}_{\text{block}}] = \sigma_E^2 + 28 \sigma_{\text{block}}^2$, giving the unbiased ANOVA estimator $\hat{\sigma}_{\text{block}}^2 = (\text{MS}_{\text{block}} - \text{MS}_E)/28$. Substituting $\hat{\sigma}_{\text{block}}^2$ and $\hat{\sigma}_E^2 = \text{MS}_E$ into V_{conf} yields the linear combination of independent mean squares required by Satterthwaite's synthesis:

$$\hat{V}_{\text{conf}} = \text{MS}_E \left(\frac{1}{m} + h_0 \right) + \frac{\text{MS}_{\text{block}} - \text{MS}_E}{28} \left(\frac{1}{m} + \frac{1}{5} \right) = c_1 \text{MS}_E + c_2 \text{MS}_{\text{block}} \quad (16)$$

where $c_2 = \frac{1/m+1/5}{28}$ and $c_1 = (1/m + h_0) - c_2$. Propagating variance-component degrees of freedom via Satterthwaite’s approximation yields effective degrees of freedom:

$$\nu_{\text{eff}} = \frac{\hat{V}_{\text{conf}}^2}{\frac{(c_1 \text{MS}_E)^2}{121} + \frac{(c_2 \text{MS}_{\text{block}})^2}{4}} \quad (17)$$

For $\mathbf{x}_{\text{MO}}^*$ ($h_0 = 0.1043$), Satterthwaite effective degrees of freedom evaluate to $\nu_{\text{eff}} = 14.0$ ($t_{0.025,14.0} = 2.145$), widening the 95% prediction interval to $[0.4677, 0.4932]$ (compared to $[0.4687, 0.4922]$ under the unadjusted $t_{0.025,121}$ approximation). For the depth-7 single-objective coordinate $\mathbf{x}_{\text{SO}}^*$ ($h_0 = 0.1210$), $\nu_{\text{eff}} = 15.1$ ($t_{0.025,15.1} = 2.131$), producing $[0.4354, 0.4612]$. We note that Satterthwaite prediction intervals pool structural lack of fit (10 df) and treatment $\times$ block interaction (96 df) into MS_E , assuming homogeneous residual variance across design cells without safeguarding against local polynomial bias. Table 10 compares predictions against empirical results.

Configuration / Response Metric	Surrogate Pred ($\hat{y}$)	95% Pred Interval (PI)	Empirical Mean $\pm$ SD	Confirmation Status
DOE Multi-Objective Optimum $\mathbf{x}_{\text{MO}}^*$ (Depth 4)				
Validation RMSE (Y_1)	0.4804	$[0.4677, 0.4932]$	0.4853 ± 0.0115	Pass (Inside 95% PI)
Holdout Test RMSE	—	—	0.4888 ± 0.0032	Holdout Test Set
Inference Latency (μs)	138.1	$[122.8, 153.4]$	142.3 ± 7.8	Pass (Inside 95% PI)
DOE Single-Objective Candidate $\mathbf{x}_{\text{SO}}^*$ (Depth 7)				
Validation RMSE (Y_1)	0.4483	$[0.4354, 0.4612]$	0.4700 ± 0.0093	Not Confirmed (Optimism: +0.0217)
Holdout Test RMSE	—	—	0.4691 ± 0.0034	Holdout Test Set
Inference Latency (μs)	187.0	$[171.1, 202.9]$	171.1 ± 10.9	Borderline (At Lower PI)

Table 10: Model Prediction vs. Empirical Confirmation Trials ($m = 10$ Fresh Seeds)

For the multi-objective compromise coordinate $\mathbf{x}_{\text{MO}}^*$ (depth 4), empirical validation RMSE of 0.4853 ± 0.0115 falls directly inside the Satterthwaite prediction interval ($[0.4677, 0.4932]$), providing pointwise confirmation consistent with the surrogate. Empirical latency was $142.33 \pm 7.79 \mu\text{s}$, also falling within $[122.8, 153.4] \mu\text{s}$. Test RMSE on the untouched holdout set ($N = 4,128$) averaged 0.4888 ± 0.0032 . Pointwise confirmation at $\mathbf{x}_{\text{MO}}^*$ does not imply global surrogate accuracy across the entire hypercube.

For the single-objective candidate $\mathbf{x}_{\text{SO}}^*$ (depth 7), the surrogate predicted $\hat{y} = 0.4483$ with Satterthwaite 95% PI $[0.4354, 0.4612]$, while empirical confirmation yielded 0.4700 ± 0.0093 . Because empirical validation RMSE lies 0.0088 above the upper bound of the Satterthwaite interval, the prediction was not confirmed. The surrogate exhibited structural optimism of +0.0217 ($6.8 \times \text{SE}(\hat{y})$), indicating that quadratic interpolation between sampled design depths $\{3, 6, 9\}$ does not adequately track the diminishing returns of deeper trees. Measured inference latency ($171.1 \pm 10.9 \mu\text{s}$) sits directly on the lower prediction interval bound ($171.1 \mu\text{s}$), representing a borderline confirmation. On the holdout test set, depth 7 achieves an external test RMSE of 0.4691 ± 0.0034 , competitive with Bayesian baselines.

8.2 Benchmark Evaluation Against Random Search and Optuna TPE

We evaluated four baseline configurations against the DOE recommended optima under equal-budget conditions (140 function evaluations):

1. **Unguided Random Search:** 20 independent runs of 140 evaluations.
2. **Single-Objective TPE:** 20 independent runs of 140 evaluations optimizing validation RMSE.
3. **Constrained TPE:** Single-objective TPE with a latency constraint ($\leq 145 \mu s$) evaluated via an analytical quadratic latency proxy model ($\tau_{\text{proxy}} = 115 + 3d + 0.8d^2 \mu s$) during trial sampling.
4. **Multi-Objective TPE:** Optimizes validation RMSE and latency simultaneously via Derringer-Suich desirability.

For each method, the representative median incumbent hyperparameter vector across the 20 search runs was retrained and evaluated across 20 fresh seeds (seeds 2001–2020, $\mathcal{S}_{\text{eval}}$) on the fixed holdout test set with pinned CPU core affinity and interleaved repeats. Table 11 summarizes the benchmark comparison.

Optimization Method	Val RMSE (95% CI)	Test RMSE (95% CI)	Pred Lat (μs)	Inplace (μs)	Depth	Search Basis
DOE Multi-Obj ($\mathbf{x}_{\text{MO}}^*$)	0.4821 [0.4786, 0.4855]	0.4884 [0.4869, 0.4899]	151.5	105.4	4	Second-Order RSM
DOE Single-Obj ($\mathbf{x}_{\text{SO}}^*, d=7$)	0.4687 [0.4654, 0.4720]	0.4690 [0.4678, 0.4702]	176.8	137.3	7	Second-Order RSM
Random Search	0.4668 [0.4636, 0.4700]	0.4675 [0.4661, 0.4690]	203.9	164.6	8	Black-Box Oracle
Bayesian TPE (SO)	0.4668 [0.4631, 0.4705]	0.4668 [0.4650, 0.4685]	206.9	167.7	8	Black-Box Oracle
Constrained TPE	0.4796 [0.4760, 0.4832]	0.4835 [0.4820, 0.4849]	152.4	112.2	4	Black-Box Oracle
Multi-Obj TPE	0.4974 [0.4940, 0.5008]	0.5032 [0.5014, 0.5051]	147.9	100.7	3	Black-Box Oracle

Table 11: Empirical Benchmark Comparison across 20 Fresh Seeds on Fixed Holdout Test Set. For each method, the representative median incumbent hyperparameter vector across the 20 search runs is evaluated across retraining seeds $\mathcal{S}_{\text{eval}}$.

Figure 5 illustrates optimization efficiency across training evaluations.




Figure 5: Sample Efficiency Convergence: Cumulative best validation RMSE vs. evaluation budget. Curves for Random Search and TPE show median trajectories $\pm$ IQR shaded bands across 20 independent runs. The DOE curve shows cumulative best in actual randomized run execution order across the pre-planned blocked design. Inset highlights fine-grained convergence between runs 20 and 140.

The empirical benchmarks reveal four primary findings:

1. **Accuracy and Latency Trade-Offs against Multi-Objective TPE:** Comparing multi-objective solutions, the DOE compromise coordinate $\mathbf{x}_{\text{MO}}^*$ achieves a test RMSE of 0.4884 at $151.5 \mu\text{s}$ latency (inplace $105.4 \mu\text{s}$), whereas Multi-Objective TPE converges to 0.5032 at $147.9 \mu\text{s}$ (inplace $100.7 \mu\text{s}$). While MO-TPE achieves lower latency by $3.6 \mu\text{s}$ (2.4%), DOE achieves lower test error (paired $t = -17.99, p = 2.16 \times 10^{-13}$; Wilcoxon $W = 0.0, p = 1.91 \times 10^{-6}$). Relative to reference point $[0.60, 250.0 \mu\text{s}]$, the point hypervolume dominated by the selected DOE compromise $\mathbf{x}_{\text{MO}}^*$ alone is $HV = 10.99$ versus $HV = 9.88$ for MO-TPE’s selected configuration, an 11.2% **point hypervolume increase**. Considering the two-point DOE set ($\mathbf{x}_{\text{MO}}^*$ and $\mathbf{x}_{\text{SO}}^*$) expands hypervolume to 12.41 (+25.7%, though evaluated across separate models rather than a joint Pareto set).
2. **Performance Relative to Constrained TPE:** When constrained during search ($\leq 145 \mu\text{s}$), Constrained TPE converges to a depth 4 configuration achieving 0.4835 test RMSE at $152.4 \mu\text{s}$. While beating DOE $\mathbf{x}_{\text{MO}}^*$ by 0.0049 RMSE, Constrained TPE exhibits a slight violation of its search-time threshold during dedicated interleaved benchmarking ($152.4 \mu\text{s}$ vs. $145.0 \mu\text{s}$), reflecting a combination of analytical proxy underestimation and hardware execution jitter between search-time estimation and dedicated measurement.
3. **Latency Advantage Over Single-Objective Baselines:** Single-objective TPE converges to depth 8, achieving 0.4668 test RMSE with an inference latency of $206.9 \mu\text{s}$. TPE is therefore 36.6% **slower** than DOE $\mathbf{x}_{\text{MO}}^*$ ($151.5 \mu\text{s}$), or equivalently, DOE latency is 26.8% **lower**.
4. **Single-Objective Competitiveness:** In pure single-objective mode, the empirical DOE depth 7 configuration ($\mathbf{x}_{\text{SO}}^*$) achieves a test RMSE of 0.4690 at $176.8 \mu\text{s}$ predict ($137.3 \mu\text{s}$ inplace), matching TPE within 0.0022 RMSE while running 14.5% faster.

9 Discussion and Limitations

The empirical findings point to several practical takeaways and methodological limits:

1. **Curvature as an Early Diagnostic:** Phase 1’s curvature test ($F = 105432.31$ against center pure error, $F \approx 4,140$ against overall pure error) detected non-planarity within 100 runs. Factorial screening provides a test for curvature before allocating evaluations to full quadratic modeling.
2. **Stochastic Seed Blocking:** Random seed blocking captured approximately 40.7% of unexplained residual variance ($\sigma_{\text{block}} \approx 0.0077$ RMSE). While small relative to dominant main effects, this variance is on par with differences separating competing algorithms, showing that blocking helps avoid selecting configurations biased by partition noise.
3. **Super-Linear Latency Dynamics:** Latency growth across discrete depths follows a quadratic trajectory ($R^2 = 0.859$, reaching $R^2 = 0.882$ under the full second-order model). Comparing standard and `inplace_predict` isolates a 39–47 μs Python and scikit-learn wrapper overhead, while base traversal time accelerates at greater depths.
4. **Structural Lack of Fit and Model Optimism:** Second-order polynomials cannot fully capture the asymptotic convergence of gradient boosting loss. This structural limitation led to a persistent RMS misfit per degree of freedom (0.0297 overall, 0.0127 on the restricted domain; unweighted cell RMS error 0.0080) and caused surrogate optimism at intermediate depth 7 (bias = +0.0217). Furthermore, Satterthwaite prediction intervals pool structural lack of fit into MS_E , which widens intervals under unexplained variance but assumes homogeneous dispersion without safeguarding against localized polynomial bias. Practitioners using RSM for machine learning should treat polynomial surrogates as local guidance maps rather than globally precise response estimators.
5. **Evaluation Budget Asymmetry and Proxy Dynamics:** The CCD evaluated 25 unique geometric configurations across 5 seed blocks (140 total runs), whereas Optuna TPE and Random Search evaluated up to 140 distinct configurations. While the 25-point structured grid enabled formal variance decomposition and hypothesis testing, adaptive Bayesian methods can explore a broader variety of interior coordinates. Conversely, Constrained TPE relied on an analytical quadratic latency proxy during search rather than online benchmarking, which contributed to slight constraint threshold overshoots during final evaluation.
6. **Scope and Operational Context:** These conclusions are established for XGBoost with 100 boosting trees on the California Housing dataset. In higher dimensions ($D > 10$), fractional factorial screening or hybrid Bayesian-DOE strategies are warranted.

10 Reproducibility Statement

All experimental code, configuration manifests, raw execution logs, and LaTeX source files are publicly available under open-source version control at:

<https://github.com/ChristosHussein/xgboost-doe-rsm> (release tag v1.0.0, pull request #1)

Experiments were executed on an AMD Ryzen 5 9600X 6-Core processor (12 logical CPUs) under Windows-11-10.0.26300-SP0 using Python 3.11.9, XGBoost 3.2.0, Scikit-Learn 1.8.0, Statsmodels 0.15.0, Scipy 1.17.1, and Optuna 5.0.0. To aid reproducibility, execution scripts pin CPU core affinity (CPU 0) and fix test partitions. Latency measurements nonetheless remain subject to operating system scheduling and hardware state.

11 Conclusion

This study evaluated classical Design of Experiments (DOE) principles for hyperparameter optimization in gradient boosted decision trees. Through sequential factorial screening, central composite design augmentation, canonical spectral analysis, and desirability optimization, the framework isolated stochastic seed variance ($\sigma_{\text{block}} \approx 0.0077$), tested response curvature, and identified a multi-objective operating coordinate ($\mathbf{x}_{\text{MO}}^*$) balancing predictive accuracy with inference latency. Empirical confirmation across fresh seeds confirmed surrogate prediction intervals at $\mathbf{x}_{\text{MO}}^*$, while identifying structural polynomial optimism at the single-objective candidate $\mathbf{x}_{\text{SO}}^*$ (+0.0217 RMSE bias). On an untouched holdout test set, DOE $\mathbf{x}_{\text{MO}}^*$ achieved lower test error at comparable latency to Multi-Objective TPE (+11.2% point hypervolume increase), while depth 7 matched single-objective Bayesian search within 0.0022 RMSE with a 14.5% latency reduction.

References

- [1] D. C. Montgomery, *Design and Analysis of Experiments*, 9th ed. Hoboken, NJ: John Wiley & Sons, 2017.
- [2] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [3] J. Bergstra and Y. Bengio, “Random search for hyper-parameter optimization,” *Journal of Machine Learning Research*, vol. 13, pp. 281–305, 2012.
- [4] J. Bergstra, R. Bardenet, Y. Bengio, and B. Kégl, “Algorithms for hyper-parameter optimization,” in *Advances in Neural Information Processing Systems*, vol. 24, 2011, pp. 2546–2554.
- [5] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, “Optuna: A next-generation hyper-parameter optimization framework,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2019, pp. 2623–2631.
- [6] G. Derringer and R. Suich, “Simultaneous optimization of several response variables,” *Journal of Quality Technology*, vol. 12, no. 4, pp. 214–219, 1980.
- [7] G. E. P. Box and K. B. Wilson, “On the experimental attainment of optimum conditions,” *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 13, no. 1, pp. 1–45, 1951.
- [8] R. H. Myers, D. C. Montgomery, and C. M. Anderson-Cook, *Response Surface Methodology: Process and Product Optimization Using Designed Experiments*, 4th ed. Hoboken, NJ: John Wiley & Sons, 2016.