

Sequential Response Surface Methodology and Central Composite Design for Multi-Objective Hyperparameter Optimization in Gradient Boosted Trees under Stochastic Nuisance Blocking

Design of Experiments & Machine Learning Research Group
Empirical Systems Engineering Laboratory

October 8, 2026

Abstract

Hyperparameter optimization (HPO) in gradient boosted decision trees is predominantly conducted using stochastic search heuristics such as uniform random sampling or Bayesian optimization via Tree-structured Parzen Estimators (TPE). While effective at unconstrained objective minimization, heuristic search strategies offer limited structural interpretability, treat stochastic nuisance variation from data partitioning as negligible noise, and lack formal statistical criteria for local curvature and model adequacy. In this investigation, we implement a sequential Design of Experiments (DOE) framework grounded in the classical methodology of Douglas C. Montgomery. We formulate a multi-objective HPO pipeline optimizing an XGBoost regressor (100 trees) on the California Housing dataset across four continuous factors: learning rate (η), maximum tree depth, subsampling fraction, and L_2 regularization (λ). Stochastic variation is isolated via a Randomized Complete Block Design (RCBD) across five deterministic seed blocks, where block-to-block data partition differences account for an estimated $\sigma_{\text{block}} \approx 0.0077$ RMSE in unexplained residual variance ($\text{ICC} = 0.4069, F = 20.21, p < 0.0001$). In Phase 1, a 2^4 full factorial design augmented with center points ($N = 100$) reveals substantial response surface curvature ($F = 105432.31, p < 10^{-15}$ against center pure error; $F \approx 4,140$ against overall pure error), decisively rejecting first-order planarity. In Phase 2, the design is augmented into a Face-Centered Central Composite Design (FCCD, $N = 140$), decomposing residual error into pure replicate error (111 df) and lack of fit (10 df). Formal lack-of-fit testing reveals significant second-order deficiency ($F = 62.21, p = 7.78e - 41$) with an RMS misfit of 0.0297. While dropping the low- η boundary ($x_1 \geq -0.5$) mechanically reduces lack-of-fit degrees of freedom and lowers RMS misfit to 0.0127, significant lack of fit remains ($F = 10.11, p = 0.0001$), highlighting the structural difficulty of fitting asymptotic boosting curves with second-order polynomials. Residual diagnostics show consistency with normality ($W = 0.9971, p = 0.9947$) and block variance equality ($W = 0.1754, p = 0.9507$), while directional variance compression across the factor domain is accommodated via HC3 robust standard errors. Phase 3 canonical spectral analysis indicates that the unconstrained stationary point lies outside the operational domain at coded distance 17.16, with wild bootstrap analysis confirming a stationary ridge system. Constrained optimization over the operational hypercube yields a predicted optimum at depth 7 ($\hat{y} = 0.4483$). In Phase 4, Derringer-Suich multi-objective desirability optimization produces an operational compromise coordinate ($\mathbf{x}^*$, depth 4, $\eta = 0.2324$) balancing predictive error with inference latency. In Phase 5, empirical confirmation across $m = 10$ fresh seeds validates surrogate predictions at $\mathbf{x}^*$ ($Y_1: 0.4853 \pm 0.0115$ inside $[0.4687, 0.4922]$; $Y_2: 142.33 \pm 7.79 \mu\text{s}$ inside $[122.82, 153.40]$). At depth 7, empirical validation RMSE (0.4700 ± 0.0093) falls above the 95% prediction interval ($[0.4363, 0.4603]$) due to structural surrogate optimism ($+0.0217$), though the coordinate itself successfully locates competitive predictive accuracy. Benchmark comparisons across 20 fresh seeds on a fixed external holdout test set ($N = 4,128$) demonstrate that DOE $\mathbf{x}^*$ delivers superior predictive accuracy at equivalent latency compared to Multi-Objective

TPE (0.4884 vs. 0.5032 test RMSE at $151.5\mu\text{s}$ vs. $147.9\mu\text{s}$, expanding single-compromise hypervolume by +11.2% from 9.88 to 10.99). Compared to single-objective TPE ($206.9\mu\text{s}$), DOE $\mathbf{x}^*$ reduces latency by 26.8% ($151.5\mu\text{s}$, making TPE 36.6% slower). Furthermore, the confirmed DOE depth 7 configuration achieves a test RMSE of 0.4690, closely matching TPE (0.4668) while running 14.5% faster.

Keywords: Design of Experiments (DOE), Response Surface Methodology (RSM), Central Composite Design (CCD), Gradient Boosted Trees, XGBoost, Multi-Objective Optimization, Randomized Complete Block Design (RCBD), Canonical Analysis, Model Adequacy.

1 Introduction and Theoretical Framework

Hyperparameter optimization (HPO) is an essential component of modern machine learning workflows. For tabular datasets, Gradient Boosted Decision Trees (GBDT), particularly as implemented in XGBoost [2], frequently achieve state-of-the-art predictive performance. However, tuning GBDT hyperparameters involves complex non-linear interactions, non-convex loss surfaces, and operational trade-offs between predictive accuracy and single-sample inference latency.

Standard industrial practice relies heavily on heuristic search algorithms, ranging from uniform random search [3] to sequential Bayesian optimization methods such as Tree-structured Parzen Estimators (TPE) [4, 5] or Gaussian Process regression. Despite their widespread adoption, these black-box optimizers present three engineering limitations:

- (i) **Absence of Parametric Attribution:** Black-box optimizers treat the model as an opaque evaluation oracle. They cannot quantify the percentage of variance explained by main effects versus two-factor interactions, nor do they provide an analytical surrogate for sensitivity analysis.
- (ii) **Confounding of Stochastic Nuisance:** Variation arising from random data partitioning and stochastic subsampling is routinely confounded with genuine hyperparameter sensitivity unless structured blocking is introduced.
- (iii) **Lack of Formal Curvature and Adequacy Tests:** Heuristic search strategies lack hypothesis-testing frameworks to formally detect whether an operational region is planar, quadratic, or characterized by structural lack of fit.

To address these limitations within an interpretable framework, this study implements a sequential Design of Experiments (DOE) pipeline based on Chapters 5, 9, 10, and 14 of Douglas C. Montgomery’s *Design and Analysis of Experiments* [1], pioneering response surface methodology originating from Box and Wilson [7] and extended by Myers et al. [8]. By coupling randomized complete block screening, curvature testing, Central Composite Design (CCD) augmentation, spectral canonical analysis, and Derringer-Suich desirability functions [6], we evaluate the extent to which classical response surface methodology can optimize GBDTs.

Every metric, sum of squares, test statistic, and diagnostic reported in this paper is derived from genuine model fits on the California Housing benchmark ($N = 20,640$). To guarantee absolute isolation from test data leakage, an external holdout test set of 4,128 observations (20%) was carved out prior to any experimental design or model training. The remaining 16,512 observations formed the development pool, which each seed block partitioned into 12,384 training (75%) and 4,128 validation (25%) samples.

2 Experimental Architecture and Mathematical Formulation

2.1 Operational Factors and Standard Coding

We define four continuous hyperparameter factors controlling structural complexity and regularization for an XGBoost regressor with $n_{\text{estimators}} = 100$:

- **Factor A** (x_1): Learning rate ($\eta \in [0.01, 0.30]$), coded on a logarithmic scale: $\xi_1 = \ln(\eta) \in [\ln(0.01), \ln(0.30)] = [-4.6052, -1.2040]$.
- **Factor B** (x_2): Maximum tree depth ($\text{depth} \in [3, 9]$), controlling structural model complexity.
- **Factor C** (x_3): Subsample fraction ($\text{subsample} \in [0.50, 1.00]$), introducing stochastic row subsampling.
- **Factor D** (x_4): L_2 leaf regularization ($\lambda \in [0.10, 10.00]$), coded on a logarithmic scale: $\xi_4 = \ln(\lambda) \in [\ln(0.10), \ln(10.00)] = [-2.3026, 2.3026]$.

Each natural variable ξ_i is mapped to dimensionless coded space $x_i \in [-1, +1]$ via:

$$x_i = \frac{\xi_i - \frac{\xi_{i,\max} + \xi_{i,\min}}{2}}{\frac{\xi_{i,\max} - \xi_{i,\min}}{2}}, \quad i \in \{1, 2, 3, 4\} \quad (1)$$

The operational domain forms a 4-dimensional hypercube $\mathcal{D} = [-1, +1]^4$. Table 1 summarizes factor coding.

Factor	Natural Variable	Transform	Low −1	Center 0	High +1
Factor A (x_1)	Learning Rate (η)	$\ln(\eta)$	0.0100	0.0548	0.3000
Factor B (x_2)	Max Tree Depth	Linear	3	6	9
Factor C (x_3)	Subsample Fraction	Linear	0.5000	0.7500	1.0000
Factor D (x_4)	L_2 Reg (λ)	$\ln(\lambda)$	0.1000	1.0000	10.0000

Table 1: Hyperparameter Factors, Operational Domains, and Design Levels

2.2 Randomized Complete Block Design (RCBD) for Stochastic Nuisance

Stochastic variation in data partitioning and subsampling introduces non-trivial variance across evaluation trials. We structure the experimental protocol as a Randomized Complete Block Design (RCBD) across $b = 5$ distinct seed blocks: $\mathcal{S} = \{42, 101, 202, 303, 404\}$. To guarantee statistical rigor, the exact assignment of split seeds and model seeds is structured as shown in Table 2.

Phase / Design Point	Point IDs	Block	Split Seed	Model Seed	Reps	Isolated Variance
Phase 1: Corner Runs	1...16	$b \in \{1..5\}$	S_b	S_b	1	Factorial + Block
Phase 1: Center Points	17	$b \in \{1..5\}$	S_b	$S_b + 1000r$	4	Subsampling Error (15 df)
Phase 2: Axial Runs	18...25	$b \in \{1..5\}$	S_b	S_b	1	Quadratic + Block

Table 2: Structure of Blocks, Split Seeds, and Model Seeds Across Experimental Phases

Specifically, each block b defines a distinct 75/25 train/validation partition of the 16,512 development observations via split seed S_b . For corner runs (1...16) and axial runs (18...25), the internal XGBoost subsampling seed is also set to S_b . For the four center point replicates within block b , the train/validation split is kept identical at S_b , but the model subsampling seed is varied via $S_b + r \times 1000$ ($r \in \{0, 1, 2, 3\}$). This distinction cleanly isolates within-block stochastic subsampling pure error (15 degrees of freedom) from across-block data-split variance (4 degrees of freedom).

The statistical model for response y_{uj} at design point u in block j is:

$$y_{uj} = \mu + \tau_u + \beta_j + \epsilon_{uj}, \quad j \in \{1, \dots, 5\} \quad (2)$$

where τ_u represents the treatment effect of design coordinate $\mathbf{x}_u$, β_j represents the block effect, and ϵ_{uj} is experimental error.

To quantify the magnitude of variance absorbed by blocking, we compute the Intraclass Correlation Coefficient (ICC) under both ANOVA and REML formulations:

$$\text{ICC} = \frac{\sigma_{\text{block}}^2}{\sigma_{\text{block}}^2 + \sigma_{\epsilon}^2} \quad (3)$$

For validation RMSE, the ANOVA estimator yields $\text{ICC} = 0.4069$ ($F = 20.21, p < 0.0001$), which matches the restricted maximum likelihood estimate ($\text{ICC}_{\text{REML}} = 0.4069$). Importantly, this 40.7% represents a share of *unexplained residual variance* rather than total experimental variation (where primary factors account for $> 90\%$ of total sum of squares). In absolute terms, the estimated block standard deviation is $\sigma_{\text{block}} = \sqrt{\sigma_{\text{block}}^2} \approx 0.0077$ RMSE. This nuisance variation is directly comparable in magnitude to the fine-grained performance differences separating competitive hyperparameter configurations (often 0.005–0.015 RMSE), underscoring the necessity of blocking.

Notably, while the block effect is non-significant in Phase 1 ($F = 0.33, p = 0.8550$), it emerges as highly significant in the Phase 2 full second-order model ($F = 20.21, p < 10^{-11}$). This discrepancy exemplifies textbook statistical shielding: in Phase 1, unmodeled quadratic curvature pooled into the residual denominator ($\text{MS}_E = 0.003185$), artificially inflating error variance and diluting the block F -ratio. Once second-order terms absorbed the curvature in Phase 2 (MS_E reduced to 0.000086), the genuine seed partition variance became cleanly identifiable. For inference latency, the block effect is negligible ($\text{ICC} = 0.0000, F = 0.75, p = 0.5591$), as execution speed is governed by hardware architecture and tree topology rather than data splits.

2.3 Dual Responses and Latency Measurement Protocol

For every trial, two responses are measured:

1. **Response Y_1 (Validation RMSE):** Root mean squared error evaluated on the 25% validation partition ($N_{\text{val}} = 4,128$) of the development pool. The holdout 20% test partition ($N_{\text{test}} = 4,128$) remains strictly untouched until final generalization testing.
2. **Response Y_2 (Single-Sample Inference Latency):** Execution time in microseconds per single sample ($\mu\text{s}/\text{sample}$).

To eliminate session drift, CPU frequency fluctuations, and OS scheduling noise, all latency measurements follow a unified measurement protocol:

- (i) **Pinned CPU Affinity:** The execution process is pinned to CPU Core 0 via Windows API `SetProcessAffinityMask(GetCurrentProcess(), 1)`.
- (ii) **Warmup and Single-Threaded Execution:** 100 untimed warmup iterations precede timed execution, with XGBoost configured for single-threaded inference (`n_jobs=1`).
- (iii) **Repeated Timed Evaluations:** Inference latency is timed across 1,000 single-sample evaluations using high-resolution monotonic clocks (`time.perf_counter_ns`), recording both the standard `predict` (with DMatrix conversion) and `inplace_predict`.
- (iv) **Interleaved Benchmark Runs:** During comparative benchmarking, evaluation runs across all competing algorithms are interleaved within the same pinned process session to guarantee an identical hardware operating state.

3 Phase 1: Screening and Curvature Assessment

Phase 1 constructs a 2^4 full factorial design consisting of $n_F = 16$ corner runs, augmented with $n_C = 4$ center points per block ($x_1 = x_2 = x_3 = x_4 = 0$). Evaluated across all 5 seed blocks, Phase 1 comprises $N_1 = 100$ genuine training runs.

We fit a first-order response surface model with all two-factor interaction terms:

$$Y = \beta_0 + \sum_{i=1}^4 \beta_i x_i + \sum_{i < j} \beta_{ij} x_i x_j + \sum_{b=2}^5 \gamma_b Z_b + \epsilon \quad (4)$$

where Z_b are dummy indicators for blocks 2 through 5. Table 3 details the Type III Sum of Squares ANOVA for Response Y_1 .

Source of Variation	SS	DF	MS	F-Stat	p value	Partial η^2
Intercept	7.309127	1	7.309127	2097.15	$< 10^{-15}$	0.9610
Block Effect $C(\text{block})$	0.004662	4	0.001166	0.33	0.8541	0.0155
Factor A: x_1 ($\ln(\eta)$)	1.245625	1	1.245625	357.40	$< 10^{-15}$	0.8079
Factor B: x_2 (Depth)	0.082155	1	0.082155	23.57	0.0000	0.2171
Factor C: x_3 (Subsample)	0.002406	1	0.002406	0.69	0.4084	0.0081
Factor D: x_4 ($\ln(\lambda)$)	0.000089	1	0.000089	0.03	0.8731	0.0003
Interaction $x_1 \cdot x_2$	0.091725	1	0.091725	26.32	0.0000	0.2364
Interaction $x_1 \cdot x_3$	0.001878	1	0.001878	0.54	0.4649	0.0063
Interaction $x_1 \cdot x_4$	0.007452	1	0.007452	2.14	0.1474	0.0245
Interaction $x_2 \cdot x_3$	0.001106	1	0.001106	0.32	0.5747	0.0037
Interaction $x_2 \cdot x_4$	0.000092	1	0.000092	0.03	0.8714	0.0003
Interaction $x_3 \cdot x_4$	0.000081	1	0.000081	0.02	0.8790	0.0003
Residual Error	0.296247	85	0.003485	—	—	—

Table 3: Phase 1 ANOVA Table (Type III Sum of Squares) for Validation RMSE (Y_1)

Factor A ($\ln(\eta)$) dominates predictive performance ($F = 357.40$, partial $\eta^2 = 0.8079$), followed by Factor B (tree depth, $F = 23.57$) and the interaction $x_1 \cdot x_2$ ($F = 26.32$, $p < 10^{-6}$).

3.1 Formal Curvature Hypothesis Test

Following Montgomery (Chapter 10), curvature is tested by comparing the average response of factorial points ($\bar{y}_F$) against center points ($\bar{y}_C$):

$$SS_{\text{Curvature}} = \frac{n_F n_C (\bar{y}_F - \bar{y}_C)^2}{n_F + n_C} \quad (5)$$

Across the 100 runs evaluated in Phase 1:

$$n_F = 80, \quad \bar{y}_F = 0.6333, \quad n_C = 20, \quad \bar{y}_C = 0.4989, \quad \bar{y}_F - \bar{y}_C = 0.1343 \quad (6)$$

yielding $SS_{\text{Curvature}} = 0.288786$. Estimating pure error variance from center point replicates within blocks (subsampling pure error, $df = 15$) yields $MS_{\text{PE, Center}} = 2.74e - 06$. The resulting test statistic is:

$$F_{\text{Curvature}} = \frac{MS_{\text{Curvature}}}{MS_{\text{PE, Center}}} = 105432.31 \quad (p < 10^{-15}) \quad (7)$$

If evaluated against overall pure error pooled across all replicated design coordinates ($df = 83$, $MS_{\text{PE, All}} = 6.98 \times 10^{-5}$), the F -statistic remains overwhelmingly significant ($F \approx 4,140$, $p < 10^{-15}$). Under either reference, the null hypothesis of planarity is rejected decisively, confirming the presence of substantial quadratic effects and necessitating second-order augmentation.

4 Phase 2: Central Composite Design Augmentation

4.1 Face-Centered Central Composite Design (FCCD)

To fit a second-order polynomial without discarding Phase 1 runs, the design is augmented into a Face-Centered Central Composite Design (FCCD, $\alpha = 1.0$), adding $2k = 8$ axial runs evaluated across the 5 seed blocks (40 additional runs, cumulative $N = 140$).

The second-order response surface model is:

$$Y = \beta_0 + \sum_{i=1}^4 \beta_i x_i + \sum_{i=1}^4 \beta_{ii} x_i^2 + \sum_{i < j} \beta_{ij} x_i x_j + \sum_{b=2}^5 \gamma_b Z_b + \epsilon \quad (8)$$

Table 4 presents the second-order ANOVA for Validation RMSE ($R^2 = 0.9951$, $\text{Adj } R^2 = 0.9944$). To guard against potential heteroscedasticity, both classical OLS test statistics and HC3 heteroscedasticity-consistent standard errors and p-values are reported.

Source of Variation	SS	DF	MS	F-Stat	OLS p	HC3 SE	HC3 p
Intercept	4.718852	1	4.718852	54749.42	$< 10^{-15}$	0.0021	$< 10^{-15}$
Block Effect $C(\text{block})$	0.006966	4	0.001742	20.21	< 0.0001	—	—
Factor A: x_1 ($\ln(\eta)$)	1.388673	1	1.388673	16111.76	$< 10^{-15}$	0.0012	$< 10^{-15}$
Factor B: x_2 (Depth)	0.101706	1	0.101706	1180.02	$< 10^{-15}$	0.0013	$< 10^{-15}$
Factor C: x_3 (Subsample)	0.002013	1	0.002013	23.36	< 0.0001	0.0012	0.0001
Factor D: x_4 ($\ln(\lambda)$)	0.000044	1	0.000044	0.52	0.4743	0.0012	0.5483
Quadratic x_1^2	0.146265	1	0.146265	1697.01	$< 10^{-15}$	0.0023	$< 10^{-15}$
Quadratic x_2^2	0.008779	1	0.008779	101.85	$< 10^{-15}$	0.0035	< 0.0001
Quadratic x_3^2	0.000005	1	0.000005	0.05	0.8188	0.0023	0.8005
Quadratic x_4^2	0.000006	1	0.000006	0.08	0.7844	0.0020	0.7197
Interaction $x_1 \cdot x_2$	0.091725	1	0.091725	1064.22	$< 10^{-15}$	0.0013	$< 10^{-15}$
Interaction $x_1 \cdot x_3$	0.001878	1	0.001878	21.79	< 0.0001	0.0013	0.0003
Interaction $x_1 \cdot x_4$	0.007452	1	0.007452	86.46	$< 10^{-15}$	0.0013	< 0.0001
Interaction $x_2 \cdot x_3$	0.001106	1	0.001106	12.83	0.0005	0.0013	0.0050
Interaction $x_2 \cdot x_4$	0.000092	1	0.000092	1.07	0.3040	0.0013	0.4118
Interaction $x_3 \cdot x_4$	0.000081	1	0.000081	0.94	0.3334	0.0013	0.4400
Residual Error	0.010429	121	0.000086	—	—	—	—

Table 4: Full Second-Order ANOVA Table for Validation RMSE (Y_1) with HC3 Robust Standard Errors

The ANOVA confirms substantial quadratic curvature along x_1^2 ($F = 1,697.01, p < 10^{-15}$; HC3 $p < 10^{-15}$) and x_2^2 ($F = 101.85, p < 10^{-15}$; HC3 $p < 0.0001$). Under HC3 robust covariance estimation, all key structural effects ($x_1, x_2, x_1^2, x_2^2, x_1 x_2, x_1 x_4$) remain statistically significant at $p < 0.0001$. Table 5 provides the corresponding ANOVA for Inference Latency (Y_2).

Source of Variation	SS	DF	MS	F-Stat	<i>p</i> value
Intercept	555450.04	1	555450.04	1903.08	$< 10^{-15}$
Block Effect $C(\text{block})$	876.94	4	219.24	0.75	0.5591
Factor A: x_1 ($\ln(\eta)$)	1265.46	1	1265.46	4.34	0.0394
Factor B: x_2 (Depth)	236841.33	1	236841.33	811.46	$< 10^{-15}$
Factor C: x_3 (Subsample)	179.04	1	179.04	0.61	0.4350
Factor D: x_4 ($\ln(\lambda)$)	483.26	1	483.26	1.66	0.2006
Quadratic x_1^2	1576.98	1	1576.98	5.40	0.0218
Quadratic x_2^2	11634.52	1	11634.52	39.86	< 0.0001
Quadratic x_3^2	8.35	1	8.35	0.03	0.8660
Quadratic x_4^2	116.71	1	116.71	0.40	0.5284
Interaction $x_1 \cdot x_2$	662.90	1	662.90	2.27	0.1344
Interaction $x_1 \cdot x_3$	14.00	1	14.00	0.05	0.8270
Interaction $x_1 \cdot x_4$	829.24	1	829.24	2.84	0.0945
Interaction $x_2 \cdot x_3$	813.68	1	813.68	2.79	0.0976
Interaction $x_2 \cdot x_4$	47.57	1	47.57	0.16	0.6871
Interaction $x_3 \cdot x_4$	52.53	1	52.53	0.18	0.6721
Residual Error	35316.19	121	291.87	—	—

Table 5: Full Second-Order ANOVA Table for Inference Latency (Y_2 , $\mu\text{s}/\text{sample}$)

4.2 Lack-of-Fit Decomposition and Causal Analysis

Following Montgomery (Section 9.3), when replicate observations exist at identical factor coordinates across blocks, the residual sum of squares with 121 df can be partitioned into Pure Error and Lack of Fit. Using nested model comparison between the second-order model (19 parameters) and a saturated cell-means model ($m = 25$ unique design points + 4 block parameters = 29 parameters):

$$\text{df}_{\text{LoF}} = 29 - 19 = 10, \quad \text{df}_{\text{PE}} = 140 - 29 = 111 \quad (9)$$

Table 6 summarizes the lack-of-fit decomposition for both responses across the full domain and within a restricted factor domain.

Model Specification	SS Resid	SS PE	SS LoF	F_{LoF}	<i>p</i> value	RMS Misfit
Y_1 : Val RMSE (Full Domain)	0.010429	0.001579	0.008850	62.21	$< 10^{-15}$	0.0297
Y_1 : Val RMSE (Domain $x_1 \geq -0.5$)	0.001528	0.001203	0.000324	10.11	0.0001	0.0127
Y_2 : Latency (μs)	35316.19	31993.82	3322.37	1.15	0.3305	18.23

Table 6: Nested Model Lack-of-Fit Decomposition and Domain Restriction Test

For Validation RMSE (Y_1) on the full operational domain, lack of fit is statistically significant ($F = 62.21, p = 7.78e - 41$). The root-mean-square misfit is $\sqrt{\text{MS}_{\text{LoF}}} = 0.0297$, which is approximately $7.8\times$ the pure error standard deviation ($\sqrt{\text{MS}_{\text{PE}}} = 0.0038$).

To inspect where this lack of fit originates across the design, we examined per-design-point residual contributions $n_i(\bar{y}_i - \hat{y}_i)^2$. The largest contributions concentrate at design coordinates where $x_1 = -1$ ($\eta = 0.01$), particularly at corners $(x_1, x_2, x_3, x_4) = (1, 1, -1, -1)$ and $(-1, -1, 1, 1)$, reflecting severe underfitting when 100 boosting iterations are paired with very low learning rates. When refitting on the restricted sub-domain $x_1 \geq -0.5$ ($\eta \geq 0.033$), the sum of squares of lack of fit drops by 96.3% (from 0.008850 to 0.000324), and the RMS misfit declines by 57% to 0.0127. However, we emphasize that this restriction drops the $x_1 = -1$ coordinate entirely, leaving only two factor levels ($x_1 \in \{0, 1\}$) and reducing df_{LoF} from 10 to 2. Consequently, a second-order polynomial in x_1 interpolates the remaining points with fewer degrees of freedom to detect higher-order curvature. Even within this restricted region, lack of fit remains statistically significant ($F = 10.11, p = 0.0001$), with RMS misfit (0.0127) remaining roughly $3\times$ the pure error standard deviation (0.0038). This analysis demonstrates that while the severe underfitting boundary at $\eta = 0.01$ drives the majority of the discrepancy, a second-order polynomial remains an approximation with non-negligible structural bias when modeling boosting dynamics.

For Latency (Y_2), lack of fit on the full CCD is not statistically significant at $\alpha = 0.05$ ($F = 1.15, p = 0.3305$). To characterize the functional form of latency growth, five candidate models were compared against empirical measurements, as reported in Table 7. For the log-linear specification, AIC and BIC values are adjusted via the Jacobian of transformation ($+2 \sum \ln(y_i) = 1,450.8$) to ensure rigorous likelihood comparability on the natural microsecond scale.

Candidate Latency Model	R^2	Adj R^2	AIC (Adjusted)	BIC (Adjusted)	RMSE (μs)
Linear (depth)	0.792	0.790	1255.2	1261.1	21.26
Quadratic (depth + depth ²)	0.859	0.857	1202.6	1211.4	17.56
Log-Linear ($\ln(\text{latency})$ depth)	0.817	0.815	1218.1	1224.0	19.57
Multi-Factor (depth, depth ² , eta, sub, lam)	0.868	0.863	1199.5	1217.1	17.18
Full CCD Second-Order (All factors + Block)	0.882	0.864	1209.6	1265.5	17.08

Table 7: Comparison of Candidate Regression Models for Inference Latency (Y_2)

The quadratic latency model ($R^2 = 0.859$, adjusted AIC = 1202.6) strongly outperforms the linear specification ($R^2 = 0.792$, AIC = 1255.2) and the log-linear specification (adjusted AIC = 1218.1), corroborating the significant x_2^2 term in ANOVA ($F = 39.86, p < 0.0001$). Latency exhibits super-linear growth across the empirical depth range ($d \in \{3, 6, 9\}$), with the quadratic fit capturing accelerating tree traversal costs while avoiding the severe over-prediction of exponential formulations. Furthermore, comparing standard prediction against `inplace_predict` reveals an average execution gap of 39–47 μs , confirming that memory transformation into XGBoost internal structures imposes a modest overhead, while base inference execution time is driven by tree structure evaluation.

5 Statistical Residual Diagnostics

To evaluate regression assumptions, diagnostic checks were conducted on externally studentized residuals:

- (a) **Normality:** Evaluated via Shapiro-Wilk on studentized residuals: $W = 0.9971, p = 0.9947$. It fails to reject normality at $\alpha = 0.05$.

- (b) **Variance Homogeneity across Blocks:** Assessed via Levene’s test across the 5 seed blocks: $W = 0.1754, p = 0.9507$. Brown-Forsythe testing across the 25 design points yields $W = 0.6872, p = 0.8552$, confirming uniform replicate dispersion across design coordinates.
- (c) **Reconciliation of Breusch-Pagan and Brown-Forsythe:** The Breusch-Pagan test detects continuous heteroscedasticity across the four-factor hypercube ($LM = 70.10, p < 0.0001$). This does not contradict the Brown-Forsythe block homogeneity: while variance across seed blocks is uniform, residual variance compresses directionally toward regions of higher η and deeper trees where the boosting loss reaches an asymptote. This continuous heteroscedasticity justifies our deployment of HC3 robust inference in Table 4.
- (d) **Independence and Run Order:** Durbin-Watson statistic evaluated in actual randomized execution order yields $DW = 1.9918$. The Ljung-Box test ($Q = 4.01, p = 0.5486$) and Wald-Wolfowitz Runs test ($p = 0.8576$) confirm the absence of serial correlation or thermal execution drift.
- (e) **Influence and Leverage:** Maximum Cook’s distance is $D_{\max} = 0.097$, well below the critical threshold $D = 1.0$.

Figure 1 displays the 4-in-1 residual diagnostic panel.




Figure 1: Residual Diagnostics Panel for Second-Order Validation RMSE Model: (a) Studentized residuals vs. fitted values; (b) Normal Q-Q probability plot ($p = 0.9947$); (c) Residuals vs. randomized run execution order ($DW = 1.9918$, Ljung-Box $p = 0.5486$, Runs test $p = 0.8576$); (d) Cook’s distance influence plot bounded well within $D = 1.0$ ($D_{\max} = 0.097$).

6 Phase 3: Canonical and Ridge Analysis

6.1 Matrix Representation and Stationary Point

The fitted second-order model is expressed in matrix notation:

$$\hat{y} = b_0 + \mathbf{x}^T \mathbf{b} + \mathbf{x}^T \mathbf{B} \mathbf{x} \quad (10)$$

where $b_0 = 0.4994$ (averaging block effects equally), $\mathbf{b} = [-0.1242, -0.0336, -0.0047, -0.0007]^T$, and:

$$\mathbf{B} = \begin{bmatrix} 0.1065 & 0.0169 & -0.0024 & -0.0048 \\ 0.0169 & 0.0261 & -0.0019 & -0.0005 \\ -0.0024 & -0.0019 & 0.0006 & 0.0005 \\ -0.0048 & -0.0005 & 0.0005 & 0.0007 \end{bmatrix} \quad (11)$$

The unconstrained stationary point $\mathbf{x}_0$ satisfies $\mathbf{b} + 2\mathbf{B}\mathbf{x}_0 = \mathbf{0}$:

$$\mathbf{x}_0 = -\frac{1}{2}\mathbf{B}^{-1}\mathbf{b} = [0.4171, 1.3524, 15.6944, -6.8009]^T \quad (12)$$

The stationary coordinate along x_3 (15.6944) lies far outside the operational cube $[-1, +1]^4$, with a Euclidean coded distance of $\|\mathbf{x}_0\| = 17.16$. The unconstrained predicted value $\hat{y}_0 = 0.4161$ represents an extrapolation artifact.

6.2 Spectral Decomposition and Ridge Analysis

Eigenvalues of $\mathbf{B}$ are $\lambda_1 = 0.000057, \lambda_2 = 0.000873, \lambda_3 = 0.022759, \lambda_4 = 0.110196$ ($\sum \lambda_i = \text{trace}(\mathbf{B}) = 0.133886$). While point estimates are positive, Rademacher wild bootstrap (2,000 replications) reveals that the 95% bootstrap interval for λ_1 spans $[-0.0039, 0.0022]$, with 68.2% of bootstrap samples yielding $\min(\lambda) \leq 0$. The system is classified as a **stationary or rising ridge**, with curvature along x_3 and x_4 indistinguishable from zero.

Constrained optimization over the hypercube $[-1, 1]^4$ with integer depth $d \in \{3, \dots, 9\}$ yields the results in Table 8.

Depth	Coded x_1	Coded x_2	Coded x_3	Coded x_4	Optimal η	Predicted RMSE	SE ($\hat{y}$)
3	0.8103	-1.0000	1.00	1.00	0.2173	0.4909	0.0032
4	0.7573	-0.6667	1.00	1.00	0.1985	0.4724	0.0030
5	0.7043	-0.3333	1.00	1.00	0.1814	0.4592	0.0032
6	0.6513	+0.0000	1.00	1.00	0.1658	0.4511	0.0033
7	0.5983	+0.3333	1.00	1.00	0.1515	0.4483	0.0032
8	0.5453	+0.6667	1.00	1.00	0.1384	0.4506	0.0031
9	0.4923	+1.0000	1.00	1.00	0.1265	0.4582	0.0034

Table 8: Best Predicted Validation RMSE at Each Discrete Tree Depth

The constrained single-objective optimum within the hypercube occurs at depth 7 ($x_1 = 0.5983, x_2 = 0.3333, x_3 = 1.0000, x_4 = 1.0000$) with predicted RMSE 0.4483. Figures 2 and 3 display the response surfaces.




Figure 2: Response Surface Topography for Validation RMSE (Y_1): (a) 2D contour slice across Factor A ($\ln(\eta)$) and Factor B (depth), passing through the constrained optimum ($x_3 = 1, x_4 = 1$). Arrow indicates trajectory toward unconstrained stationary point outside domain; (b) 3D wireframe projection.




Figure 3: Inference Latency Modelling: (a) Empirical mean latency with 95% CI error bars vs. linear ($R^2 = 0.792$) and quadratic ($R^2 = 0.859$) fits. Note the substantial fixed overhead ($\approx 105 \mu s$) arising from Python execution and tree traversal; (b) Distribution boxplots across discrete depths.

7 Phase 4: Multi-Objective Desirability Optimization

Minimizing predictive error requires deeper trees, which increases inference latency. To identify a principled compromise, we employ Derringer-Suich desirability functions [6]:

$$d_r(Y_r) = \begin{cases} 1.0, & Y_r < L_r \\ \left(\frac{U_r - Y_r}{U_r - L_r} \right)^{s_r}, & L_r \leq Y_r \leq U_r \\ 0.0, & Y_r > U_r \end{cases} \quad (13)$$

The composite desirability D is formulated with explicit importance weights (w_1, w_2):

$$D = (d_1(Y_1)^{w_1} \cdot d_2(Y_2)^{w_2})^{\frac{1}{w_1 + w_2}} \quad (14)$$

In the baseline scenario, equal weights $w_1 = 1, w_2 = 1$ are utilized with engineering limits $L_1 = 0.450, U_1 = 0.700$ for RMSE and $L_2 = 100.0, U_2 = 180.0 \mu\text{s}$ for latency ($s_1 = s_2 = 1.0$).

Maximizing D within the operational domain with an interior safety buffer against domain-boundary saturation ($\eta \leq 0.25$) yields the recommended multi-objective compromise operating coordinate:

$$\begin{aligned} \mathbf{x}^* &= [0.8500, -0.6667, 1.0000, -0.0812]^T \\ \implies & (\eta = 0.2324, \text{depth} = 4, \text{subsample} = 1.0000, \lambda = 0.8295) \end{aligned} \quad (15)$$

achieving $D = 0.6782$ ($d_1 = 0.8783, d_2 = 0.5236$). Predicted metrics at $\mathbf{x}^*$ are $\hat{Y}_1 = 0.4804$ and $\hat{Y}_2 = 138.11 \mu\text{s}$.

Table 9 presents a sensitivity analysis across alternative specification limits and weightings. In unconstrained boundary grid search, several scenarios push η directly to the upper domain limit $\eta = 0.3000$ ($x_1 = 1.0$, reaching $D = 0.6952$). This occurs because with a fixed budget of 100 boosting trees, the algorithm aggressively favors larger learning rates to accelerate residual minimization. We selected $\mathbf{x}^*$ at $\eta = 0.2324$ ($x_1 = 0.8500$) as the operational benchmark coordinate to prevent boundary pinning and avoid boundary extrapolation instability during empirical confirmation.

Scenario	RMSE Bounds [L_1, U_1]	Latency Bounds [L_2, U_2]	Weights (w_1, w_2)	Opt Depth	Opt η	Best D
Standard Config	[0.45, 0.70]	[100, 180]	(1, 1)	4	0.3000	0.6952
Strict Latency	[0.45, 0.70]	[90, 140]	(1, 2)	3	0.3000	0.2206
Strict Accuracy	[0.44, 0.60]	[100, 200]	(2, 1)	5	0.2135	0.7293
Narrow Range	[0.46, 0.65]	[110, 160]	(1, 1)	4	0.3000	0.6597

Table 9: Desirability Optimization Sensitivity Analysis Across Bounds and Weights

Figure 4 illustrates the empirical Pareto frontier and the location of $\mathbf{x}^*$.




Figure 4: Multi-Objective Pareto Trade-off between Validation RMSE and Inference Latency. Scatter points from CCD design are colored by tree depth using *cividis*. The empirical Pareto frontier is traced in red, with DOE multi-objective optimum $\mathbf{x}^*$ (depth 4), DOE single-objective optimum (depth 7), confirmation means ($n = 10$), and benchmark incumbents highlighted.

8 Phase 5: Empirical Confirmation and Benchmark Comparison

8.1 Confirmation Experiments at Multi-Objective and Single-Objective Optima

To test surrogate fidelity, $m = 10$ confirmation trials were executed at $\mathbf{x}^*$ across fresh seeds disjoint from the design ($\mathcal{S}_{\text{conf}} = \{505, \dots, 1414\}$). To account for the fact that fresh seeds carry block variance, the 95% prediction interval for the mean of m future trials incorporates both parameter covariance and block variance:

$$\text{PI}_{0.95}(\bar{y}_m) = \hat{y} \pm t_{0.025, 121} \sqrt{\text{MS}_E \left(\frac{1}{m} + \mathbf{x}^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{x} \right) + \sigma_{\text{block}}^2 \left(\frac{1}{m} + \frac{1}{5} \right)} \quad (16)$$

In addition, we tested the single-objective optimum (depth 7) to directly evaluate model bias against empirical performance. Table 10 compares predictions against empirical results.

Configuration / Response Metric	Surrogate Pred ($\hat{y}$)	95% Pred Interval (PI)	Empirical Mean $\pm$ SD	Confirmation Status
DOE Multi-Objective Optimum x^* (Depth 4)				
Validation RMSE (Y_1)	0.4804	[0.4687, 0.4922]	0.4853 ± 0.0115	Pass (Inside 95% PI)
Holdout Test RMSE	—	—	0.4888 ± 0.0032	Holdout Test Set
Inference Latency (μs)	138.1	[122.8, 153.4]	142.3 ± 7.8	Pass (Inside 95% PI)
DOE Single-Objective Optimum (Depth 7)				
Validation RMSE (Y_1)	0.4483	[0.4363, 0.4603]	0.4700 ± 0.0093	Empirical Opt (Bias: +0.0217)
Holdout Test RMSE	—	—	0.4691 ± 0.0034	Holdout Test Set
Inference Latency (μs)	187.0	[171.1, 202.9]	171.1 ± 10.9	Borderline (At Lower PI)

Table 10: Model Prediction vs. Empirical Confirmation Trials ($m = 10$ Fresh Seeds)

For the multi-objective optimum $\mathbf{x}^*$ (depth 4), the empirical validation RMSE of 0.4853 ± 0.0115 falls directly inside the 95% prediction interval ([0.4687, 0.4922]), confirming surrogate validity. Empirical latency was $142.33 \pm 7.79 \mu s$, also falling well within [122.8, 153.4] μs . Generalization test RMSE evaluated on the untouched external holdout set ($N = 4, 128$) averaged 0.4888 ± 0.0032 .

For the single-objective optimum (depth 7), the surrogate predicted $\hat{y} = 0.4483$ with 95% PI [0.4363, 0.4603], whereas empirical confirmation yielded 0.4700 ± 0.0093 . Because empirical RMSE lies 0.0097 above the upper bound of the prediction interval, the formal hypothesis of numerical prediction fidelity is rejected. The surrogate exhibited structural optimism of $\text{bias} = +0.0217$ ($6.8 \times \text{SE}(\hat{y})$). This structural misfit arises because depth 7 lies between the sampled design levels $\{3, 6, 9\}$, where the quadratic interpolant in x_2 underestimates the flattening of validation loss gains. Simultaneously, measured inference latency ($171.1 \pm 10.9 \mu s$) sits directly on the lower prediction interval bound (171.1 μs), representing a borderline confirmation. Nonetheless, despite numerical optimism, evaluating depth 7 empirically achieves an external test RMSE of 0.4691 ± 0.0034 , successfully capturing competitive performance.

8.2 Benchmark Evaluation Against Random Search and Optuna TPE

We evaluated four baseline configurations against the DOE recommended optima under equal-budget conditions (140 function evaluations):

1. **Unguided Random Search:** 20 independent runs of 140 evaluations.
2. **Single-Objective TPE:** 20 independent runs of 140 evaluations optimizing validation RMSE.
3. **Constrained TPE:** Single-objective TPE with latency constraint ($\leq 145 \mu s$) during search.
4. **Multi-Objective TPE:** Optimizes validation RMSE and latency simultaneously.

Final incumbents from each method were evaluated across 20 fresh seeds (seeds 2001–2020, $\mathcal{S}_{\text{eval}}$) on the fixed holdout test set with pinned CPU core affinity and interleaved repeats. Table 11 summarizes the benchmark comparison.

Optimization Method	Val RMSE (95% CI)	Test RMSE (95% CI)	Pred Lat (μ s)	Inplace (μ s)	Depth	Search Basis
DOE Multi-Obj (x^*)	0.4821 [0.4786, 0.4855]	0.4884 [0.4869, 0.4899]	151.5	105.4	4	Second-Order RSM
DOE Single-Obj ($d=7$)	0.4687 [0.4654, 0.4720]	0.4690 [0.4678, 0.4702]	176.8	137.3	7	Second-Order RSM
Random Search	0.4668 [0.4636, 0.4700]	0.4675 [0.4661, 0.4690]	203.9	164.6	8	Black-Box Oracle
Bayesian TPE (SO)	0.4668 [0.4631, 0.4705]	0.4668 [0.4650, 0.4685]	206.9	167.7	8	Black-Box Oracle
Constrained TPE	0.4796 [0.4760, 0.4832]	0.4835 [0.4820, 0.4849]	152.4	112.2	4	Black-Box Oracle
Multi-Obj TPE	0.4974 [0.4940, 0.5008]	0.5032 [0.5014, 0.5051]	147.9	100.7	3	Black-Box Oracle

Table 11: Empirical Benchmark Comparison across 20 Fresh Seeds on Fixed Holdout Test Set

Figure 5 illustrates optimization efficiency across training evaluations.




Figure 5: Sample Efficiency Convergence: Cumulative best validation RMSE vs. evaluation budget. Curves for Random Search and TPE show median trajectories $\pm$ IQR shaded bands across 20 independent runs. The sequential DOE curve shows cumulative best in actual randomized run execution order. Inset highlights fine-grained convergence between runs 20 and 140.

The empirical benchmarks reveal four primary findings:

1. **DOE Delivers Superior Accuracy at Equivalent Latency over Multi-Objective TPE:** Comparing multi-objective solutions, the DOE optimum $\mathbf{x}^*$ achieves a test RMSE of 0.4884 at 151.5 μ s latency (inplace 105.4 μ s), whereas Multi-Objective TPE converges to 0.5032 at 147.9 μ s (inplace 100.7 μ s). While MO-TPE is marginally faster by 3.6 μ s (2.4%), DOE achieves substantially lower test error (paired $t = -17.99, p = 2.16 \times 10^{-13}$;

Wilcoxon $W = 0.0, p = 1.91 \times 10^{-6}$). Comparing hypervolume (HV) relative to reference point $[0.60, 250.0 \mu\text{s}]$, DOE $\mathbf{x}^*$ alone achieves $HV = 10.99$ versus $HV = 9.88$ for MO-TPE—an 11.2% **hypervolume expansion**. When the two-point DOE set ($\mathbf{x}^*$ and depth 7) is considered, hypervolume expands to 12.41 (+25.7%).

2. **Performance Relative to Constrained TPE:** When constrained during search ($\leq 145 \mu\text{s}$), Constrained TPE converges to a depth 4 configuration achieving 0.4835 test RMSE at $152.4 \mu\text{s}$. While beating DOE $\mathbf{x}^*$ by 0.0049 RMSE, Constrained TPE exhibits a slight violation of its search-time threshold during dedicated interleaved benchmarking ($152.4 \mu\text{s}$ vs. $145 \mu\text{s}$), reflecting hardware execution variability between search-time estimation and dedicated measurement.
3. **Latency Advantage Over Single-Objective Baselines:** Single-objective TPE pushes to depth 8, achieving 0.4668 test RMSE at the expense of an inference latency of $206.9 \mu\text{s}$. TPE is therefore 36.6% **slower** than DOE $\mathbf{x}^*$ ($151.5 \mu\text{s}$), or equivalently, DOE latency is 26.8% **lower**.
4. **Single-Objective Competitiveness:** In pure single-objective mode, the empirical DOE depth 7 configuration achieves a test RMSE of 0.4690 at $176.8 \mu\text{s}$ predict ($137.3 \mu\text{s}$ inplace), matching TPE within 0.0022 RMSE while running 14.5% faster.

9 Discussion and Limitations

The empirical findings highlight key engineering insights along with important methodological boundaries:

1. **Curvature as an Early Diagnostic:** Phase 1’s curvature test ($F = 105432.31$ against center pure error, $F \approx 4,140$ against overall pure error) decisively detected non-planarity within 100 runs. In production pipelines, screening designs provide formal hypothesis testing before committing resources to full second-order modeling.
2. **Stochastic Seed Blocking:** Random seed blocking captured approximately 40.7% of unexplained residual variance ($\sigma_{\text{block}} \approx 0.0077$ RMSE). While small relative to dominant main effects, this variance is on par with the fine margins separating competing algorithms, demonstrating why blocking is essential to avoid choosing sub-optimal configurations based on split artifacts.
3. **Super-Linear Latency Dynamics:** Latency growth across discrete depths follows a quadratic trajectory ($R^2 = 0.859$, reaching $R^2 = 0.882$ under the full second-order model). Comparing standard and `inplace_predict` isolates a 39–47 μs memory conversion overhead, while base traversal time accelerates at greater depths.
4. **Structural Lack of Fit and Model Optimism:** Second-order polynomials cannot fully capture the asymptotic convergence of gradient boosting loss. This structural limitation led to a persistent RMS misfit (0.0297 overall, 0.0127 on the restricted domain) and caused surrogate optimism at intermediate depth 7 (bias = +0.0217). Practitioners using RSM for machine learning should treat polynomial surrogates as local guidance maps rather than globally precise response estimators.
5. **Evaluation Budget Asymmetry:** The CCD evaluated 25 unique geometric configurations across 5 seed blocks (140 total runs), whereas Optuna TPE and Random Search evaluated up to 140 distinct configurations. While the 25-point structured grid enabled formal variance decomposition and hypothesis testing, adaptive Bayesian methods can explore a broader variety of interior coordinates.

6. **Scope and Operational Context:** These conclusions are established for XGBoost with 100 boosting trees on the California Housing dataset. In higher dimensions ($D > 10$), fractional factorial screening or hybrid Bayesian-DOE strategies are warranted.

10 Reproducibility Statement

All experimental code, configuration manifests, raw execution logs, and LaTeX source files are publicly available under open-source version control at:

<https://github.com/ChristosHussein/xgboost-doe-rsm> (commit 4528425)

Experiments were executed on an AMD Ryzen 5 9600X 6-Core processor (12 logical CPUs) under Windows-11-10.0.26300-SP0 using Python 3.11.9, XGBoost 3.2.0, Scikit-Learn 1.8.0, Statsmodels 0.15.0, Scipy 1.17.1, and Optuna 5.0.0. To maximize experimental reproducibility, code execution pins CPU core affinity (CPU 0) and fixes external test partitions. We note, however, that inference latency remains subject to operating system scheduling jitter and background hardware state.

11 Conclusion

This study demonstrated the application of classical Design of Experiments (DOE) principles to hyperparameter optimization in gradient boosted decision trees. Through sequential factorial screening, central composite design augmentation, canonical spectral analysis, and desirability optimization, the framework isolated stochastic seed variance ($\sigma_{\text{block}} \approx 0.0077$), identified significant response curvature, and located a multi-objective operating coordinate ($\mathbf{x}^*$) balancing predictive accuracy with inference latency. Empirical confirmation across fresh seeds verified surrogate prediction intervals at the multi-objective coordinate, while highlighting structural polynomial optimism at the single-objective optimum (+0.0217 RMSE bias). Benchmarks on an untouched holdout test set demonstrated that DOE delivers superior accuracy at equivalent latency compared to Multi-Objective TPE (expanding single-compromise hypervolume by +11.2%), matches single-objective Bayesian search within 0.0022 RMSE with a 14.5% latency reduction, and reduces latency by 26.8% relative to unconstrained Bayesian optimization.

References

- [1] D. C. Montgomery, *Design and Analysis of Experiments*, 9th ed. Hoboken, NJ: John Wiley & Sons, 2017.
- [2] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [3] J. Bergstra and Y. Bengio, “Random search for hyper-parameter optimization,” *Journal of Machine Learning Research*, vol. 13, pp. 281–305, 2012.
- [4] J. Bergstra, R. Bardenet, Y. Bengio, and B. Kégl, “Algorithms for hyper-parameter optimization,” in *Advances in Neural Information Processing Systems*, vol. 24, 2011, pp. 2546–2554.
- [5] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, “Optuna: A next-generation hyper-parameter optimization framework,” in *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 2019, pp. 2623–2631.

- [6] G. Derringer and R. Suich, “Simultaneous optimization of several response variables,” *Journal of Quality Technology*, vol. 12, no. 4, pp. 214–219, 1980.
- [7] G. E. P. Box and K. B. Wilson, “On the experimental attainment of optimum conditions,” *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 13, no. 1, pp. 1–45, 1951.
- [8] R. H. Myers, D. C. Montgomery, and C. M. Anderson-Cook, *Response Surface Methodology: Process and Product Optimization Using Designed Experiments*, 4th ed. Hoboken, NJ: John Wiley & Sons, 2016.